

**LEGALVISION: A KNOWLEDGE-DRIVEN AI
FRAMEWORK FOR
LEGAL UNDERSTANDING AND TRUST ASSESSMENT**

Project ID – 25-26J-127

Final Report

K. Sharan - IT22732736

S. Sivanuja - IT22348784

Wicramasinghe D A T N - IT22113504

Maxwell Y - IT22344892

BSc (Hons) degree in Information Technology
Specializing in Data Science

Department of Information Technology

Sri Lanka Institute of Information Technology
Sri Lanka

April 2026

**LEGALVISION: A KNOWLEDGE-DRIVEN AI
FRAMEWORK FOR
LEGAL UNDERSTANDING AND TRUST ASSESSMENT**

Project ID – 25-26J-127

Final Report

K. Sharan - IT22732736

S. Sivanuja - IT22348784

Wicramasinghe D A T N - IT22113504

Maxwell Y - IT22344892

BSc (Hons) degree in Information Technology
Specializing in Data Science

Department of Information Technology

Sri Lanka Institute of Information Technology
Sri Lanka

April 2026

DECLARATION

I declare that this is my own work, and this proposal does not incorporate without acknowledgement any material previously submitted for a degree or diploma in any other university or Institute of higher learning and to the best of our knowledge and belief it does not contain any material previously published or written by another person except where the acknowledgement is made in the text.

Name	Student ID	Signature
K. Sharan	IT22732736	
S. Sivanuja	IT22348784	S. Sivanuja .
Wicramasinghe D A T N	IT22113504	
Maxwell Y	IT22344892	

The supervisor/s should certify the proposal report with the following declaration.
The above candidates are carrying out research for the undergraduate Dissertation under my supervision.


.....

30/04/2026
.....

Signature of the Supervisor
(Ms. Dushanthi Kuruppu)

Date

ABSTRACT

The complexity of legal documents, particularly in the domain of property law, often presents significant challenges to both legal professionals and non-expert users. In Sri Lanka, property law documents such as deeds, rental agreements, and leases are crucial but can be difficult to interpret due to dense language, complex clauses, and legal jargon. This research presents the development of an AI-powered legal document analysis system, specifically designed to automate the review, evaluation, and interpretation of property law documents. The system aims to provide legal professionals with a tool that streamlines document review and helps non-expert users gain a better understanding of legal terms, clauses, and potential risks.

The AI system integrates advanced Natural Language Processing (NLP) techniques, such as Named Entity Recognition (NER), to identify and extract key legal entities (e.g., property owners, tenants, dates, and contract clauses) from property law documents. In addition, the system incorporates risk detection algorithms that automatically assess and flag risky clauses (such as unilateral termination clauses and unclear ownership terms). A major innovation of this system is its Bias Detection feature, which identifies biased clauses that may disproportionately favor one party over another (e.g., landlord vs. tenant) and offers a trust score to assess document reliability.

Another key component of this research is the Explainable AI (XAI) framework, which ensures that the system's decision-making process is transparent and auditable. Through the IRAC (Issue, Rule, Application, Conclusion) reasoning method, the system not only provides insights but also clearly explains the legal rationale behind its decisions, making it easier for users to understand how certain clauses are flagged. This is especially valuable for non-expert users who may not be familiar with legal terminology but still need to make informed decisions.

The system also includes a Knowledge Graph that stores interconnected legal data, such as relationships between parties, property transactions, and legal clauses,

enabling users to perform advanced queries and explore the data in a structured manner. This graph ensures that the system remains dynamic and can integrate real-time legal updates, such as new statutes, court rulings, and land registry changes.

Throughout the development of this system, the research team has demonstrated how AI, machine learning, and explainable AI (XAI) can be effectively applied to automate and enhance the accuracy of legal document analysis. The system is designed not only to assist legal professionals in streamlining their document review process but also to empower non-experts, such as property buyers, sellers, and renters, to understand complex legal agreements and make more informed decisions.

In conclusion, this AI-powered system transforms the process of legal document analysis, making it more efficient, accessible, and transparent. By automating the detection of legal risks, biases, and providing explainable insights, this system has the potential to significantly impact Sri Lankan property law, and could be expanded to other legal domains in the future, improving the accessibility and clarity of legal documents for all users.

ACKNOWLEDGEMENT

I would like to express my heartfelt gratitude to all those who have supported me throughout the course of this project and contributed to its completion.

First and foremost, I would like to thank my supervisor, Ms. Dushanthi Kuruppu, for her invaluable guidance, support, and encouragement throughout this research. Her insights, expertise, and constructive feedback have been instrumental in shaping the project.

I also extend my deepest appreciation to my co-supervisor, Ms. Adya Dissanayake, for her continued support and insightful suggestions, which greatly enhanced the quality of this work.

I would like to thank my fellow team members—S. Sivanuja, Wicramasinghe D A T N, and Maxwell Y—for their hard work, dedication, and collaboration. The success of this project is a result of the collective effort and commitment from each of us.

Additionally, I would like to acknowledge the faculty members, colleagues, and friends who have provided assistance and encouragement at various stages of the project. Their support has been crucial to completing this dissertation.

Finally, I would like to express my deepest gratitude to my family for their unwavering support and belief in me throughout this entire process. Their encouragement and understanding have been invaluable.

List Of Tables

Table 1 - Evaluation Metrics — Component 1	46
Table 2 - Classification Performance Metrics	46
Table 3 - Summarization Performance (ROUGE Scores).....	46
Table 4 - Testing Plan — Explainable Legal Reasoning.....	48
Table 5 - Test Results Summary — Data Pipeline and Dataset Quality Testing.....	49
Table 6 - Test Results Summary — Training Pipeline, Evaluation Framework, and Deployment Testing	50
Table 7 - Evaluation Metrics — Bias and Risk Evaluation Module	52
Table 8 - Expected Classification Performance Metrics — Bias and Risk Evaluation	52
Table 9 - Evaluation Metrics — Dynamic Legal Knowledge Graph Module	54
Table 10 - Expected NER Performance Metrics — Knowledge Graph Module	55
Table 11 - Overall Evaluation Summary — All Four Components	56

LIST OF ABBREVIATIONS

Abbreviation	Full Form
AI	Artificial Intelligence
NLP	Natural Language Processing
NER	Named Entity Recognition
BERT	Bidirectional Encoder Representations from Transformers
Legal-BERT	Legal-BERT
XAI	Explainable Artificial Intelligence
IRAC	Issue, Rule, Application, Conclusion
QLoRA	Quantized Low-Rank Adaptation
CoT	Chain-of-Thought
F1-Score	F1-Score
UI	User Interface
API	Application Programming Interface
OCR	Optical Character Recognition
AUC	Area Under the Curve
SQL	Structured Query Language
JSON	JavaScript Object Notation
CSV	Comma-Separated Values

1. INTRODUCTION

1.1 Background Study and Literature Review

1.1.1 Background Study

Legal documents, particularly those in property law, are inherently complex and difficult to interpret. Property deeds, leases, rental agreements, and other such legal documents often feature dense legal terminology and intricate clause structures, which pose significant challenges for non-experts to understand. This lack of accessibility is not only a barrier to the general public but also to legal professionals who, despite their expertise, need to invest substantial time and effort to analyze and ensure the accuracy of legal documents.

In Sri Lanka, property law governs all property transactions, including land purchases, leases, inheritance, and ownership transfers. However, these documents typically exist in an unstructured form (such as PDFs), making it difficult to perform systematic analysis, verify ownership chains, or check for inconsistencies and risks over time. The current system does not allow for real-time analysis, dynamic updates, or easy access to interconnected data, thus limiting its effectiveness for legal professionals, property buyers, sellers, and financial institutions.

In response to these challenges, there has been increasing interest in the application of Artificial Intelligence (AI), specifically Natural Language Processing (NLP), to automate the review and analysis of property law documents. By transforming raw legal text into structured, actionable data, AI can facilitate faster, more accurate assessments and reduce the cognitive load on users. However, existing AI tools for legal analysis are still limited in their ability to provide jurisdiction-specific insights, traceability, and transparency in decision-making. This gap is particularly evident in the context of Sri Lankan property law, which lacks dedicated, machine-readable datasets and tailored AI systems for processing local legal content.

The research aims to address this gap by developing a robust, AI-driven system for property law analysis, focusing on legal document summarization, bias and risk detection, and explainable AI-driven decision-making. This system will not only streamline the legal document review process but also make legal information more accessible, interpretable, and actionable for both legal professionals and non-expert stakeholders.

1.1.2 Literature Review

Over the past few years, there has been significant progress in the development of AI and NLP systems for legal document analysis. However, the application of these technologies in property law remains an emerging field with several challenges and opportunities. The following key areas have been explored in previous research:

1. **Legal Document Analysis and Summarization:** Early work in legal text summarization relied on extractive methods, where important sentences were directly pulled from the original document. While effective for preserving the exact wording of the legal text, extractive methods often failed to capture the broader context and relationships between clauses. More recently, abstractive summarization techniques, powered by Transformer-based models such as BERT and its legal variants like LEGAL-BERT, have shown greater promise in providing meaningful summaries that go beyond simple extraction. These models generate human-readable summaries that capture the core meaning of the document while maintaining legal accuracy.
2. **Bias Detection and Risk Assessment in Legal Documents:** Bias detection in legal documents, particularly in property law, has gained attention due to the potential for unfair or imbalanced clauses that favor one party over another. Existing tools like CLAUDETTE focus on identifying unfair clauses in consumer-facing contracts but are limited when it comes to domain-specific contracts, such as property agreements. There is a notable lack of systems designed to detect bias in property law clauses, particularly with respect to party-directional bias (i.e., whether the clause favors the landlord or tenant). This research builds upon these earlier efforts by developing a multi-label

classifier that can detect both the type of bias and the severity of risk associated with different clauses.

3. **Knowledge Graphs for Legal Intelligence:** Knowledge graphs have emerged as a powerful tool for organizing and structuring complex legal data. They allow for the integration of various legal concepts, statutes, and case law into an interconnected framework that facilitates easy querying and retrieval. However, while there is significant work on knowledge graphs for general legal domains, few systems have been tailored for property law. LegalVision's approach to creating a dynamic, auto-updating knowledge graph for Sri Lankan property law is an innovative step toward filling this gap.
4. **Explainable AI in Legal Decision-Making:** One of the most pressing issues in legal AI systems is the lack of explainability. Most AI-driven systems, especially those using black-box models like deep learning, fail to provide transparent reasoning for their outputs. This is particularly concerning in high-stakes domains like law, where decisions must be auditable and justifiable. Recent research has focused on enhancing the transparency and trustworthiness of AI models by integrating explainable AI (XAI) techniques. For example, using techniques like Chain-of-Thought reasoning and argumentation graphs, AI systems can generate step-by-step justifications for their decisions, improving both user trust and the system's practical usability.
5. **AI-based Legal Reasoning and Localized Models:** There is growing recognition of the need for jurisdiction-specific legal AI systems. Models trained on general-purpose legal datasets often fail to account for regional legal nuances, such as those found in Sri Lankan property law. The importance of adapting models to local legal systems is highlighted by recent work on domain-specific fine-tuning, such as the use of QLoRA for cost-effective domain adaptation. This research shows that jurisdiction-specific fine-tuning leads to better performance and more reliable outputs when compared to models trained on general legal corpora.

In summary, the literature reveals that while there has been substantial progress in legal document analysis, bias detection, and knowledge graph construction, there

remains a significant gap in providing a comprehensive, explainable, and jurisdiction-specific AI system for property law. This research aims to address these gaps by developing a tailored AI solution for Sri Lankan property law, combining the strengths of legal document summarization, bias and risk detection, knowledge graph integration, and explainable legal reasoning.

1.2. Research Gap

Despite significant advancements in Artificial Intelligence (AI) and Natural Language Processing (NLP) for legal document analysis, several gaps remain, particularly in the context of property law in Sri Lanka. The literature and existing systems reveal critical areas where research and technology are still lacking, and where this study aims to make a substantial contribution.

1.2.1. Limited Domain-Specific Legal Tools for Property Law

Most existing AI-powered legal systems focus on general legal domains, such as consumer contracts and commercial agreements, while neglecting the unique requirements of property law. Tools like CLAUDETTE, which focus on identifying unfair clauses in consumer-facing contracts, are not tailored for property agreements such as leases, rental contracts, and property transfers. This research highlights the absence of domain-specific datasets and AI models for Sri Lankan property law, which limits the effectiveness of existing tools in the local context. There is a critical need for AI systems designed to understand the specific terminology, risk factors, and legal nuances inherent in property-related agreements.

1.2.2. Absence of Jurisdiction-Specific Legal Models

Another significant gap is the lack of jurisdiction-specific AI models. Most legal NLP systems are trained on general-purpose corpora (e.g., U.S. or UK law), which do not account for the specific statutes, case law, or legal conventions governing property transactions in Sri Lanka. This is a crucial limitation for any system intended to support Sri Lankan legal professionals and citizens. There are few models trained

specifically on Sri Lankan property law, and even fewer that are adapted to the local legal context. The challenge lies in developing AI systems that can understand and navigate the specific legal landscape of Sri Lanka, particularly its property laws, land registration systems, and associated legal documentation.

1.2.3. Lack of Explainability in Legal AI Systems

Explainability remains a major challenge in AI systems, especially in legal domains where transparency is essential for decision-making and trust. Many existing AI tools, particularly those based on deep learning, operate as “black boxes,” providing outputs without clear reasoning. Legal professionals and non-experts alike require systems that not only provide insights but also offer step-by-step justifications for their conclusions. This is particularly true in property law, where decisions regarding ownership, encumbrances, and risks can have significant financial and legal consequences. The current lack of explainability in AI-driven legal systems limits their practical application in real-world legal scenarios. This research addresses this gap by focusing on developing explainable AI models that generate auditable, traceable, and transparent legal reasoning, particularly for property law applications.

1.2.4. Inability to Integrate Dynamic, Real-Time Data

Legal systems in Sri Lanka often rely on static, offline documents such as PDFs, which can make it challenging to access up-to-date property information and legal decisions. Current legal AI systems fail to accommodate dynamic data streams, such as new property transactions, updated statutes, or judgments. These static systems hinder real-time analysis, which is crucial for keeping legal insights relevant and actionable. There is a lack of systems capable of incorporating real-time data from sources like government gazettes, land registries, or court decisions. This research aims to fill this gap by developing a dynamic system that can continually update its legal knowledge base, providing current and accurate legal insights.

1.2.5. Insufficient Tools for Risk and Bias Detection in Property Agreements

Another gap identified in the literature is the insufficient attention paid to bias and risk detection within property law agreements. While there has been progress in detecting biased clauses in commercial contracts and consumer agreements, property-related clauses often involve complex relationships that are not easily captured by current AI systems. Property agreements frequently contain clauses that disproportionately favor one party (e.g., the landlord or the buyer), leading to legal imbalances. These biases can be difficult to detect without a domain-specific approach that understands the subtleties of property law. Furthermore, the lack of risk detection systems, such as those that assess the potential legal or financial risks in property transactions, leaves legal professionals and the public vulnerable to oversight. This research introduces a novel multi-label classification system that can detect bias and assess risk severity in property law clauses, addressing these gaps in current systems.

1.2.6. Absence of User-Friendly Legal Systems for Non-Experts

Finally, while there is considerable research into legal AI systems, few of these tools are designed with non-experts in mind. The majority of legal AI tools are built for legal professionals, requiring advanced knowledge of legal terminology and concepts to navigate effectively. This excludes large segments of the population, particularly in Sri Lanka, who may benefit from automated legal assistance but lack the necessary expertise. This research aims to bridge this gap by developing a user-friendly interface that presents legal information in a simplified, accessible manner, allowing non-experts to navigate and understand complex legal documents, such as property agreements, without needing a legal background.

Summary of the Research Gap

In conclusion, the primary research gaps identified include the lack of domain-specific AI tools for Sri Lankan property law, insufficient jurisdiction-specific models, the need for explainable AI in legal decision-making, the inability to integrate real-time data, inadequate systems for detecting bias and risk in property agreements, and the absence of user-friendly systems for non-expert users. This research aims to address

these gaps by developing a comprehensive, explainable, and domain-adapted AI system for property law in Sri Lanka.

This section identifies the critical research gaps and positions your project as a valuable contribution to the field of legal AI, particularly in the context of Sri Lankan property law. Let me know if you'd like further elaboration or adjustments!

1.3. Research Problem

The research problem addressed by this study stems from the challenges faced in understanding, interpreting, and analyzing property law documents within Sri Lanka. Legal documents, particularly property deeds, rental agreements, and other property-related contracts, are dense, complex, and often inaccessible to the general public and even to legal professionals without significant effort. These documents, typically available in unstructured formats like PDFs, contain highly technical legal terminology and intricate clause structures that make it difficult to extract meaningful, actionable insights.

The research problem can be broken down into several key challenges:

1.3.1. Lack of Accessible and Machine-Readable Legal Data

Legal documents in Sri Lanka, especially property deeds, are primarily stored in non-machine-readable formats such as PDF. This makes it nearly impossible to efficiently query or analyze legal data at scale. Without structured and machine-readable data, legal professionals and the public are forced to rely on manual reviews, which are time-consuming, error-prone, and inefficient. The absence of a comprehensive, digital system capable of extracting, organizing, and querying key information from these documents presents a significant barrier to effective property law analysis.

1.3.2. Absence of AI-Based Systems for Sri Lankan Property Law

Despite the growing interest in AI-based legal tools globally, Sri Lankan property law remains underrepresented in terms of AI applications. Existing AI models and systems in legal NLP (Natural Language Processing) are often trained on datasets from other jurisdictions and cannot fully comprehend the specific laws, practices, and terminology used in Sri Lanka's property legal framework. As a result, current AI tools cannot provide accurate or reliable analysis of property law documents in the Sri Lankan context. The need for an AI-based system capable of understanding, analyzing, and extracting insights from property documents specific to Sri Lanka is a critical research problem.

1.3.3. Insufficient Transparency and Explainability in Legal AI Systems

AI systems, particularly those based on deep learning, often act as "black boxes," offering predictions or classifications without providing clear, interpretable reasoning. This lack of explainability is a significant issue in the legal domain, where decisions must be auditable and justifiable. Legal professionals, who rely on transparency and accountability, are hesitant to trust AI outputs that are not supported by clear reasoning. For property law applications, where decisions about ownership, encumbrances, and risk can have profound financial and legal consequences, the need for an explainable AI system becomes even more critical. This research aims to address this gap by developing an AI system that not only performs legal analysis but also provides auditable, transparent, and understandable reasoning for its conclusions.

1.3.4. Difficulty in Identifying Bias and Risk in Property Law Clauses

Property law agreements often contain clauses that may be biased in favor of one party, such as a landlord or seller. These biases can lead to unfair or imbalanced legal outcomes, especially if they go unnoticed. Moreover, many property agreements also contain clauses that may pose legal or financial risks, such as unclear ownership transfers, hidden encumbrances, or unilateral termination conditions. Currently, there are no comprehensive AI systems in Sri Lanka that can automatically detect and assess bias and risk in property law clauses. This research problem centers on developing a

system that can identify biased clauses, evaluate the risk severity of different sections, and provide a trust score that highlights potential issues within property agreements.

1.3.5. Limited Real-Time Data Integration and Updates

The legal landscape is dynamic, with laws, case precedents, and property transactions constantly evolving. However, existing legal tools often rely on static data sources, which can quickly become outdated. The ability to integrate real-time data, such as updates from government gazettes, legal judgments, or land registries, is a crucial research problem that has been largely ignored in current systems. For property law applications, having an up-to-date, real-time knowledge base is essential for providing accurate and current legal advice. This research aims to develop a system that continuously updates its knowledge base, ensuring that it reflects the latest legal developments and property transactions.

1.3.6. Accessibility of Legal Information for Non-Experts

Legal information, particularly in property law, is often inaccessible to non-experts due to its technical nature and the complexity of legal terms. The general public, including property buyers, renters, and sellers, often lack the expertise needed to interpret legal documents fully. This results in individuals entering into agreements without understanding the risks or legal implications of certain clauses. There is a significant gap in the availability of tools that simplify and present legal information in a manner that is easy to understand. This research seeks to address this by developing a system that is not only capable of analyzing legal documents but also presents the results in a user-friendly format that is understandable to both legal professionals and non-experts alike.

Summary of the Research Problem

In summary, the research problem revolves around the lack of accessible, AI-driven tools tailored to Sri Lankan property law, the absence of transparency and explainability in existing legal AI systems, and the failure to address bias and risk

detection within property agreements. Furthermore, the static nature of current legal systems and the difficulty in making legal information accessible to non-experts add to the problem. This research aims to bridge these gaps by developing an explainable, AI-powered system capable of analyzing property law documents, detecting bias and risk, and presenting the results in an understandable and actionable format, while integrating real-time data for continuous updates.

1.4. Research Objectives

1.4.1. Main Objectives

The main objective of this research is to develop an AI-driven system for legal document analysis and decision-making, specifically tailored for Sri Lankan property law. This system will leverage Natural Language Processing (NLP) and explainable AI (XAI) techniques to analyze property-related legal documents, such as deeds, leases, and rental agreements. The key focus areas include:

- Automating the extraction and interpretation of property law documents: The system will extract critical legal information from property agreements and transform unstructured document data (e.g., PDFs, DOCX) into structured, machine-readable formats.
- Identifying and assessing risk and bias: The system will automatically detect biased clauses and evaluate the risk severity (e.g., financial imbalance, unilateral termination, etc.) within property law agreements.
- Providing explainable insights: The research aims to provide transparent, auditable, and understandable reasoning behind the AI-generated conclusions. This includes the integration of Chain-of-Thought reasoning models to explain how specific clauses or decisions were made.
- Real-time updates and dynamic data integration: The system will continuously incorporate updates from various legal and property-related sources (e.g., land registries, government gazettes, and case law), ensuring that its legal knowledge base remains current.

1.4.2. Specific Objectives

To achieve the main objectives outlined above, the specific objectives of this research include:

1. Develop a domain-specific Named Entity Recognition (NER) model:
 - Train and fine-tune an NER model that is tailored for Sri Lankan property law, capable of extracting relevant entities (e.g., parties, property details, dates, etc.) from unstructured legal documents.
 - Achieve a high F1-score to ensure the NER model's accuracy in recognizing critical legal entities within property documents.
2. Design and implement a multi-label classifier:
 - Build a classifier that can identify different types of bias (e.g., landlord-favoring clauses) and evaluate the risk severity (e.g., high, moderate, or low) of individual clauses.
 - Develop a trust score system that quantifies the overall bias and risk of the entire document, providing a single actionable score for legal professionals and the public.
3. Integrate a Knowledge Graph:
 - Create a property law knowledge graph that maps the relationships between legal entities, statutes, property documents, and case law. This will support efficient querying, exploration, and analysis of property law-related data.
 - Use the knowledge graph to link legal clauses to relevant statutes, case precedents, and other legal documents, enabling enhanced legal reasoning and decision support.
4. Ensure explainability and transparency in decision-making:
 - Implement explainable AI techniques, such as Chain-of-Thought reasoning and argumentation graphs, to provide step-by-step explanations for the AI's legal conclusions and decisions.
 - Design the system to highlight and justify detected bias or risk, making the AI-generated outputs more trustworthy and auditable.

5. Develop a user-friendly front-end interface:
 - Create a web-based user interface using React that allows users (legal professionals and non-experts) to upload legal documents, view detected risks and biases, and receive detailed explanations of the analysis.
 - Ensure the interface is intuitive and provides interactive visualizations (e.g., risk distribution graphs, timeline views) to simplify the understanding of complex legal data.
6. Real-time data integration:
 - Develop mechanisms to integrate real-time legal data, including updated statutes, judgments, and property records, into the system, ensuring that the analysis remains accurate and up-to-date.

1.4.3. Business Objectives

In addition to the technical and academic objectives, the research also aims to meet certain business-related objectives that will make the system practical and valuable for real-world users, particularly in the context of Sri Lankan property law. These business objectives include:

1. Market Readiness and Commercial Viability:
 - Ensure that the developed system is commercially viable and ready for adoption by law firms, property professionals, financial institutions, and government agencies. The system should significantly reduce the time spent on document analysis, thus providing a cost-effective solution for property law professionals.
2. User Adoption and Engagement:
 - Promote widespread adoption of the system among non-experts, including property buyers, renters, and sellers, by focusing on ease of use and providing clear, actionable insights. The system will empower individuals without legal expertise to make more informed decisions regarding property transactions.
3. Support for Legal Transparency and Fairness:

- Through the identification of biased clauses and risks in property documents, the system will contribute to legal transparency and fairness in property transactions. It can help mitigate the risk of unfair agreements by flagging potentially problematic clauses for review, thereby enhancing legal protections for the general public.
4. Scalability and Future Integration:
- Ensure that the system can be scaled and adapted for use in other legal domains, such as family law, contract law, or corporate law, after successful deployment in property law. This will enhance the overall utility and market potential of the system in various legal sectors.
5. Compliance with Legal and Data Privacy Regulations:
- Ensure that the system complies with data privacy laws (e.g., GDPR, local Sri Lankan data protection regulations) and legal standards to maintain the trust of users and ensure that sensitive legal information is processed securely.

Summary of Research Objectives

In summary, this research aims to develop an AI-powered system that automates the extraction and analysis of Sri Lankan property law documents, providing risk and bias detection, real-time updates, and explainable legal reasoning. The specific objectives target domain-specific NLP techniques, the creation of a knowledge graph, and the development of a user-friendly interface, all while addressing business goals like market readiness, user adoption, and legal transparency.

2. Methodology

2.1 Introduction

The methodology of this research outlines the systematic process followed to develop an AI-powered legal document analysis system specifically for Sri Lankan property law. The purpose of this system is to automate the extraction of meaningful legal insights from property-related documents, such as deeds, leases, and rental agreements, while also identifying biases and assessing risks within these documents. By utilizing Natural Language Processing (NLP) and Explainable AI (XAI), this system will not only provide automated document analysis but will also ensure that users can trust and understand the AI's decision-making process. This methodology ensures the development of a system that is effective, scalable, and user-friendly.

To achieve this, the research methodology follows a structured, multi-phase approach that is designed to progressively build the system and refine it through continuous testing and iteration. The process begins with research design, where the overall framework for the system is defined, followed by feasibility studies to assess the technical, operational, and economic viability of the project. This leads to the requirement gathering and analysis phase, during which input is collected from key stakeholders, including legal professionals and non-expert users, to ensure that the system meets real-world needs. The final phase is the system design, which focuses on creating a detailed blueprint for the system's architecture, models, and user interface.

At the core of this methodology is the adoption of the agile development framework, which emphasizes iterative development, constant feedback, and flexibility. This approach allows for rapid adjustments to the system's design based on user feedback and testing results. It ensures that the development of the system is not only technically sound but also aligned with the evolving needs of the users throughout the project lifecycle.

2.1.1 Research Design and Framework

The research design for this study follows a systematic approach aimed at developing an AI-powered legal document analysis system specifically tailored for Sri Lankan property law. The framework for this research is based on established principles of System Development Life Cycle (SDLC) and agile development, ensuring flexibility, iterative progress, and continuous improvement throughout the project. The design phase defines the high-level goals and the overall structure of the system, laying the foundation for its development, testing, and deployment.

The primary focus of this research is to create a multi-phase pipeline that begins with the collection and preprocessing of property law documents and ends with a fully functional user interface capable of presenting the analysis results. This design allows for the integration of several key components that drive the functionality of the system:

1. Data Collection and Preprocessing:

- The research starts by acquiring relevant property-related legal documents such as property deeds, leases, and rental agreements. These documents, which are typically in unstructured formats (e.g., PDFs), need to be processed and converted into machine-readable formats. Text extraction methods, such as Optical Character Recognition (OCR) for scanned documents and text parsing techniques for digital documents, will be employed. This is the first crucial step to make the documents amenable to legal analysis.

2. NLP Model Development:

- After preprocessing the documents, Natural Language Processing (NLP) models will be developed to process and analyze the extracted text. A Named Entity Recognition (NER) model will be trained to identify and extract essential legal entities from the documents, such as parties involved (e.g., buyer, seller), property details (e.g., address, type of property), and other relevant legal terms. Additionally, models for

bias and risk detection will be incorporated to flag potentially problematic clauses in the documents, such as those that may unfairly favor one party or pose financial or legal risks.

3. Knowledge Graph Integration:

- A significant aspect of this research is the creation of a legal knowledge graph, which will store and map relationships between legal entities, statutes, property transactions, and related case law. This will allow the system to create complex queries and visualize relationships between different legal elements. For example, if a user queries about the ownership history of a property, the knowledge graph can pull related documents, statutes, and previous transactions to provide a comprehensive answer.

4. Explainable AI (XAI):

- One of the most critical components of the research design is the integration of explainable AI (XAI) techniques. Many AI-driven systems, especially those based on deep learning, are often considered “black-box” systems, meaning they provide predictions or classifications without explaining how they arrived at their conclusions. This lack of transparency can reduce user trust, especially in legal applications. To address this, the system will employ Chain-of-Thought reasoning, which will provide step-by-step explanations for the AI's decision-making process. For example, if a clause in a property agreement is flagged as biased, the system will provide an explanation detailing why the clause was flagged, referencing specific legal criteria or relationships in the document.

5. User Interface (UI) and Experience (UX):

- The final aspect of the research design is the development of a user-friendly interface. The system will present the results of the legal document analysis through interactive visualizations and simplified insights. Legal professionals will receive detailed reports and explanations, while non-expert users (such as property buyers or tenants) will be presented with clear, actionable insights. The goal is to

design an interface that is intuitive, easy to navigate, and visually appealing, allowing users to easily interact with the system and understand its outputs.

Throughout the research process, the methodology is grounded in an iterative design framework, where each phase builds on the previous one. This allows for flexibility, enabling continuous improvements based on user feedback, performance evaluation, and testing. The agile approach ensures that the system evolves over time, incorporating new insights, technologies, and user needs, thus maintaining its relevance and utility.

By following this structured research design and framework, the system will not only provide accurate legal analysis but also ensure that the results are transparent, explainable, and accessible to a wide range of users. This approach guarantees that the final product will effectively address the challenges and gaps currently present in Sri Lankan property law analysis, particularly with regard to the need for automated, accessible, and trustworthy legal tools.

2.1.2 Feasibility Study and Planning

Before embarking on the development of the AI-powered legal document analysis system for Sri Lankan property law, a comprehensive feasibility study was conducted to assess the viability of the project from several important perspectives: technical, operational, economic, and legal. This study was crucial to ensuring that the system could be realistically developed, deployed, and sustained, while meeting the needs of both legal professionals and non-expert users. The feasibility study and planning phase allowed us to identify potential challenges, define the scope, and set a clear direction for the project's execution.

- **Technical Feasibility**

The technical feasibility focused on determining whether the necessary tools, technologies, and methodologies were available and could be effectively applied to

the development of the system. A key part of this feasibility assessment was evaluating the Natural Language Processing (NLP) models and AI technologies that would be required for the analysis of property law documents.

The system relies heavily on Named Entity Recognition (NER) for extracting key legal entities from documents, such as parties, dates, property details, and legal clauses. The feasibility study evaluated the availability and effectiveness of existing NER models, such as SpaCy, BERT, and Legal-BERT, which have been proven successful in legal text analysis in other jurisdictions. The study found that these models could be adapted to the Sri Lankan legal context, although additional domain-specific fine-tuning would be necessary to improve accuracy and relevance.

Furthermore, the integration of knowledge graphs was assessed. The feasibility of using a graph database like Neo4j to store relationships between legal entities, property documents, statutes, and case law was confirmed. The study also considered the integration of real-time data, such as property transaction records and updated legal statutes, into the knowledge graph. It was determined that existing digital platforms, such as land registries and government gazettes, could be leveraged to ensure that the system remains current with real-time updates.

- Operational Feasibility

The operational feasibility aimed to assess whether the system could be practically implemented and used by the target audience. This included legal professionals, such as lawyers and judges, as well as non-expert users, including property buyers, renters, and tenants. The system's ability to handle both complex legal analysis for professionals and simplified insights for non-experts was a key factor.

User feedback and consultation with legal professionals and property experts confirmed the need for an intuitive, user-friendly interface. For professionals, the system would need to provide detailed and accurate legal analysis, while for non-experts, it would need to present clear, actionable insights without overwhelming users with legal jargon. The interface would include features such as risk visualization, trust

scores, and step-by-step explanations of legal decisions. The study found that such a system could be developed using existing web technologies like React for the frontend and FastAPI for the backend, ensuring a seamless user experience.

Additionally, the system's ability to handle a large volume of legal documents was a key consideration. The feasibility study confirmed that the system would need to be scalable to accommodate an increasing number of documents, users, and real-time updates. The modular design of the system was found to be adaptable to future growth, allowing for the expansion of its capabilities to include other legal domains (e.g., family law, contract law) as needed.

- Economic Feasibility

The economic feasibility study evaluated the cost-effectiveness of the project, ensuring that it could be completed within budget while providing long-term value to stakeholders. The development of the AI-powered system was found to be financially viable, primarily due to the use of open-source technologies and the adoption of pre-trained models that could be fine-tuned to the specific needs of Sri Lankan property law.

For instance, tools like FastAPI, SpaCy, Neo4j, and React were selected for the backend and frontend development because they are open-source and have well-established communities, reducing the costs associated with licensing proprietary software. Additionally, using pre-trained models such as BERT allowed the project to minimize the need for expensive data collection and model training from scratch. Instead, the models could be fine-tuned using existing property law datasets and real-world data specific to Sri Lanka, significantly lowering development costs.

The system's maintenance costs were also considered, with the study identifying the potential for cloud-based deployment on platforms like Vercel and AWS to provide cost-effective scalability and maintenance. The use of cloud infrastructure would reduce the need for heavy upfront investments in hardware and make the system more

accessible to a wider range of users, including legal professionals and property stakeholders.

- Legal Feasibility

The legal feasibility of the system was a crucial aspect of the study, particularly with regard to the privacy and security of legal data. Given the sensitive nature of the data being processed, the system was designed to comply with data protection regulations, such as GDPR (General Data Protection Regulation) and Sri Lankan data protection laws. The feasibility study ensured that the system would use encryption and other security measures to safeguard users' legal documents and personal information.

Moreover, the intellectual property (IP) considerations were addressed, ensuring that the developed system would not violate any existing legal rights while ensuring the proper attribution of data sources. The research also considered how the system could support legal audits and be used in a transparent manner, providing clear and justifiable explanations for its outputs, which is essential in legal contexts.

- Planning and Project Scope

The planning phase of the project defined the scope, timeline, and key milestones for the development of the system. The project scope includes the development of the core system, which encompasses document parsing, risk and bias detection, knowledge graph construction, and the user interface. Additionally, the system's scalability and ability to update in real time were prioritized.

The timeline for the project was broken into distinct phases, with regular checkpoints for testing and evaluation. These phases included:

1. Phase 1: Data collection and preprocessing.
2. Phase 2: Development of NLP models (NER and risk detection).
3. Phase 3: Construction of the knowledge graph and integration of real-time data.
4. Phase 4: Frontend development and user interface design.

5. Phase 5: Testing, evaluation, and deployment.

By breaking the project into smaller phases, the team was able to allocate resources efficiently and ensure that the system could be tested and refined incrementally.

In conclusion, the feasibility study confirmed that the development of an AI-powered legal document analysis system for Sri Lankan property law is technically, operationally, economically, and legally feasible. By leveraging open-source technologies, utilizing domain-specific data, and following best practices for security and compliance, the project is well-positioned for successful development. The planning phase further ensures that the system will meet the needs of its intended users, providing valuable insights into property law, while being scalable and user-friendly.

2.1.3 Requirement Gathering and Analysis

The requirement gathering and analysis phase is one of the most critical stages in the system development process, as it ensures that the AI-powered legal document analysis system will meet the needs of its users and align with the overall goals of the research. During this phase, both functional and non-functional requirements were identified, and the needs of different stakeholders were carefully considered to ensure the system is effective, usable, and scalable.

Stakeholder Involvement

The first step in gathering requirements was to engage with key stakeholders, including legal professionals, property experts, and potential system users (both legal and non-experts). The primary stakeholders were:

1. **Legal Professionals:** Lawyers, judges, and legal assistants who work with property law documents regularly. Their input was critical in understanding the types of legal insights and tools they would require to perform their work more efficiently and accurately.
2. **Non-Expert Users:** These include property buyers, renters, sellers, and anyone who interacts with property documents but may not have formal legal training.

Their input focused on how the system can make complex legal content more accessible and actionable for people without a legal background.

3. **Technical Experts:** Developers and IT professionals who provided insights into the system architecture, database design, and security measures required to ensure the system's functionality and scalability.

The requirement gathering phase was carried out using interviews, surveys, and discussions with stakeholders to understand their pain points, needs, and expectations. This helped to ensure that the final system would be aligned with real-world requirements and provide value to its users.

Functional Requirements

The functional requirements outline the core features and functionalities that the system must provide. These features were derived from the needs of legal professionals, non-expert users, and the overall goals of the research:

1. **Legal Document Parsing and Extraction:**
 - The system must be capable of extracting key legal information from a wide range of property law documents, such as deeds, rental agreements, leases, and other contractual documents. This includes extracting critical entities such as parties involved (e.g., buyer, seller), property details (e.g., location, size, type), dates (e.g., signing date, contract duration), and legal clauses (e.g., termination clauses, ownership rights).
 - The extracted data must be converted into a structured machine-readable format, enabling further analysis and integration with other components like the knowledge graph.
2. **Bias and Risk Detection:**
 - The system must detect bias in legal clauses, such as clauses that disproportionately favor one party over another (e.g., landlord vs. tenant). This functionality is particularly critical in property law, where contracts often contain clauses that are unfair or imbalanced.

- The system should also assess the risk level of detected clauses, highlighting potential legal or financial risks for users. For example, clauses related to ownership transfer, liabilities, or unilateral termination should be flagged for review, with a risk score assigned based on their severity.

3. Knowledge Graph Construction and Integration:

- The system should automatically create a knowledge graph to map relationships between legal entities, such as property owners, tenants, contracts, and legal clauses. This knowledge graph will also integrate relevant statutes, case law, and previous transactions, enabling the system to provide deeper legal insights.
- The knowledge graph will be designed to support complex queries, such as tracking the ownership history of a property or identifying legal precedents that may impact the interpretation of a clause in a contract.

4. Search and Query Capabilities:

- Users must be able to perform searches within the system to retrieve relevant legal information. For example, users could query the knowledge graph for ownership records or search for specific legal clauses (e.g., those related to mortgage terms or property boundaries).
- The system should support semantic queries that understand the context of the request and provide relevant results, even if the user does not use exact legal terminology.

5. Explainable AI (XAI):

- Given the complexity of legal documents, explainability is a key requirement. The system must provide clear and understandable explanations for its analyses and outputs. For instance, if a biased clause is detected, the system should provide a step-by-step explanation detailing why the clause is considered biased, referencing specific legal criteria or contextual information that justifies the decision.
- This is essential for building trust in the system, especially for users who may not have a legal background.

Non-Functional Requirements

In addition to the core features, the system must meet several non-functional requirements, which relate to the system's performance, usability, and scalability:

1. Scalability:

- The system must be able to handle a growing number of legal documents, especially as the database expands with more property transactions and legal records over time. This includes ensuring that the system can manage increasing amounts of data and users without compromising performance.
- The system must also be scalable to accommodate additional legal domains, such as family law or contract law, in the future.

2. Performance:

- The system must process legal documents efficiently. Users should be able to upload and analyze documents quickly, even if they contain a large amount of data.
- The system should be able to run complex queries across the knowledge graph in real time, providing users with fast, relevant results.

3. Usability:

- The user interface (UI) must be intuitive and user-friendly. Legal professionals should have access to detailed analysis, while non-expert users should be able to understand the results easily. This includes visualizing the analysis in a clear, actionable manner, such as showing risk distribution graphs and trust scores.
- Ease of navigation and simplicity are key to ensuring that users, especially those without legal expertise, can interact with the system effectively.

4. Security and Privacy:

- Given that the system will handle sensitive legal data, the system must ensure data security and comply with data privacy regulations (e.g., GDPR and Sri Lankan data protection laws). Legal documents must be stored securely, and user data must be protected against unauthorized access.
- The system will incorporate encryption and secure access protocols to ensure that all legal documents and user interactions are safeguarded.

User Requirements

User requirements were gathered from consultations with legal professionals and non-expert stakeholders. The system must meet the following needs:

1. Legal Professionals:

- Legal professionals require detailed, accurate legal analysis from the system. They need the system to be capable of providing insightful, data-driven legal advice based on the documents it analyzes.
- It should help identify potential risks, such as legal ambiguities or unclear ownership chains, and assist in reviewing property law documents more efficiently.

2. Non-Expert Users:

- For non-experts, the system must simplify complex legal concepts and present them in a way that is easy to understand. This includes providing visualized risk scores, summarized findings, and actionable insights.
- The system should also help non-expert users feel confident in understanding the terms and clauses in property documents, such as whether certain clauses are risky or biased.

The requirement gathering and analysis phase has ensured that the AI-powered legal document analysis system meets both the functional and non-functional needs of its users. By understanding the challenges faced by legal professionals and non-expert users, the system will be designed to provide detailed legal analysis with explanations,

risk detection, and user-friendly insights. These requirements form the foundation for the system's development, ensuring that it is both technically capable and user-centric, ready to meet the evolving needs of Sri Lankan property law and beyond.

2.1.4 System Design

The system design phase outlines the architecture, structure, and components of the AI-powered legal document analysis system tailored for Sri Lankan property law. This phase plays a critical role in ensuring that the system is not only functional but also scalable, efficient, and user-friendly. The system design is divided into key components, which include the architectural overview, database schema design, workflow logic, and user interface design. Each component was carefully crafted to address the specific needs of legal professionals and non-expert users, while ensuring high performance and ease of use.

A. Architectural Overview

The system follows a multi-layered architecture, which ensures modularity, scalability, and flexibility. The architecture is divided into three primary layers: Presentation Layer (Frontend), Application Layer (Backend), and Data Layer (Storage). This separation of concerns allows for easier maintenance and upgrades, as well as providing a clear structure for the system's operations.

1. Frontend (Presentation Layer):

- The frontend is the user interface (UI) that interacts with the users. It is built using React, a modern JavaScript library for building interactive user interfaces. React was chosen because it provides a dynamic, responsive, and user-friendly interface for both legal professionals and non-expert users.
- This layer enables users to upload property documents, view the analysis results, interact with the knowledge graph, and visualize the detected biases, risks, and trust scores. The frontend ensures that complex legal insights are presented in a simplified and understandable

way, with features like interactive graphs, summarized reports, and explanation tooltips for detected issues.

2. Backend (Application Layer):

- The backend is developed using FastAPI, a high-performance web framework for building APIs with Python. FastAPI ensures that the backend processes requests efficiently, handling document parsing, NLP model execution, and knowledge graph queries.
- This layer is responsible for handling user requests, such as uploading documents, analyzing legal clauses, detecting bias, and querying the knowledge graph. It also manages the data processing pipeline, including the integration of real-time data and the execution of AI models like NER and bias/risk detection.
- The application layer also integrates with external APIs or databases to retrieve the latest legal data (e.g., government gazettes, statutes) and updates the knowledge graph in real-time, ensuring the system remains current with legal changes.

3. Data Layer (Storage):

- The data layer is responsible for storing structured and unstructured data, including property law documents, extracted legal entities, and knowledge graph data.
- MongoDB is used to store the raw documents and their metadata (e.g., document type, date of upload), while Neo4j is used to manage the knowledge graph, which maps the relationships between entities, such as property owners, legal clauses, statutes, and case law.
- The knowledge graph stores this data in a format that can be easily queried and explored, enabling the system to perform complex legal queries and relationship-based searches (e.g., retrieving all previous property transactions for a specific owner).

B. Database Schema Design

The database schema design ensures that the data is stored in a structured manner, facilitating easy retrieval and efficient querying. The schema is divided into two main parts: Document Data and Knowledge Graph Data.

1. Document Data:

- This includes the raw property documents (e.g., deeds, rental agreements) and their associated metadata. The schema includes fields such as:
 - Document ID: A unique identifier for each document.
 - Document Type: Type of document (e.g., deed, lease).
 - Date: The date the document was created or uploaded.
 - Trust Score: A score assigned based on the risk and bias analysis of the document.
 - Parties Involved: Information about the property owners, tenants, or other involved parties.

2. Knowledge Graph Data:

- The knowledge graph is structured to represent legal entities and their relationships. The schema consists of:
 - Node Types: Different legal entities, such as Property Owners, Legal Clauses, Property Details, and Legal Statutes.
 - Relationship Types: Connections between entities, such as ownership relations, clause references, and legal references.
 - Properties: Metadata for each node and relationship (e.g., ownership duration, clause terms, related legal case numbers).
- This structure supports graph traversal queries, such as tracing ownership history, identifying relationships between clauses, or querying applicable statutes.

C. Game Flow Logic

While the term game flow logic is traditionally associated with interactive design in games, in this context, it refers to the user interaction flow within the system. This

flow is critical for ensuring a seamless user experience as they upload documents, interact with the analysis results, and query the knowledge graph.

The workflow logic is designed as follows:

1. Document Upload: Users upload a property document (e.g., deed, lease agreement) via the user interface.
2. Document Parsing: The system extracts the text from the document, using OCR for scanned files and text parsing for digital documents.
3. Risk and Bias Detection: The system applies the NER model to identify legal entities and the bias/risk detection model to analyze clauses for fairness and risk severity.
4. Knowledge Graph Integration: The system integrates the extracted data into the knowledge graph, linking relevant legal clauses, statutes, and case law.
5. User Interaction: The analysis results are presented to the user, including a trust score, detected risks, and biases. Users can click on detected issues for detailed explanations and legal references.
6. Search and Query: Users can query the knowledge graph to explore relationships between different legal entities and retrieve relevant legal information.

The flow is designed to ensure that users are guided through a clear, intuitive process, with actionable insights at each step.

D. User Interface and Experience Design

The user interface (UI) is a crucial aspect of the system design, as it determines how users will interact with the system. The design focuses on providing an intuitive and user-friendly experience for both legal professionals and non-expert users.

1. Simplified Interactions for Non-Experts: Non-legal users (e.g., property buyers, renters) need a system that simplifies complex legal documents and provides clear, understandable insights. The interface will present key

findings, such as risk scores, biased clauses, and recommended actions, in an easy-to-understand format. Visual elements like graphs and charts will be used to display the severity of risks and biases.

2. **Detailed Analysis for Legal Professionals:** Legal professionals will have access to more detailed reports and analysis, including full-text breakdowns of detected clauses, case law references, and statutes linked to the document. The UI will allow them to drill down into the details, view explanations for detected issues, and navigate the knowledge graph to explore related documents or legal precedents.
3. **Interactive Features:** Features such as search bars for querying legal data, clickable graphs for visualizing relationships, and tooltips for providing additional context on findings will make the interface both functional and engaging.
4. **Responsive Design:** The system will be designed to function smoothly on various devices, ensuring that users can access and interact with it on desktops, tablets, or smartphones, increasing its accessibility.

Conclusion of System Design

The system design for this AI-powered legal document analysis system combines modern technologies like NLP, knowledge graphs, and explainable AI (XAI) with a modular architecture that is scalable, efficient, and user-friendly. The use of a multi-layered architecture ensures that the system is easy to maintain and update as needed, while the design of the knowledge graph and workflow logic enables the system to offer deep, meaningful insights into property law documents. The system's user interface is tailored to ensure that both legal professionals and non-expert users can interact with the system easily, empowering them to make better-informed decisions based on clear, actionable insights.

2.2 Implementation

The implementation phase is where the system begins to take shape through development and testing. In this section, we will outline the breakdown of the sprint

process, the tools and technologies used in the development, the AI model implementation, model deployment, and the final step of game integration and scoring. Each of these steps plays a critical role in bringing the AI-powered legal document analysis system to life, ensuring that the system meets its objectives of automating legal analysis, detecting bias and risks, and providing explainable insights.

2.2.1 Sprint Breakdown

The implementation process is carried out in sprints, following the Agile methodology, which divides the work into smaller, manageable chunks. Each sprint focuses on achieving specific goals and milestones, allowing for continuous progress and flexibility throughout the development process.

1. Sprint 1 – Data Collection and Preprocessing:

- During this sprint, the main goal is to collect a sufficient number of Sri Lankan property law documents (e.g., property deeds, leases) and preprocess them. Preprocessing involves extracting text from scanned PDF documents using Optical Character Recognition (OCR) and converting digital documents into a machine-readable format.
- This sprint also includes data cleaning, where irrelevant text is removed, and the documents are formatted for input into the NER (Named Entity Recognition) model.

2. Sprint 2 – NER Model Development and Training:

- This sprint focuses on developing the NER model that will extract key legal entities from the documents (e.g., parties involved, property details, legal clauses). The model is trained using labeled datasets of legal entities, focusing on the specific needs of Sri Lankan property law.
- During this sprint, model training is done iteratively, and performance is evaluated using metrics such as Precision, Recall, and F1-score to ensure the model's effectiveness in identifying entities.

3. Sprint 3 – Risk and Bias Detection Model:

- The primary focus of this sprint is to implement the risk and bias detection model. This model is trained to identify biased clauses (e.g., tenant vs. landlord bias) and assess the severity of risks in the property documents. The detection model is based on supervised learning techniques and is refined by continually training the model with labeled data from property law documents.
 - The model assigns a risk score to each document and flags any clauses that exhibit high bias or present potential legal risks.
4. Sprint 4 – Knowledge Graph Construction and Integration:
- This sprint involves developing the legal knowledge graph that will connect the extracted legal entities and relationships (e.g., ownership, property details, case law) into a structured, queryable format.
 - Neo4j, a graph database, is used to create the knowledge graph. The graph schema is designed to represent nodes for entities like property owners and legal clauses, and edges to represent relationships such as ownership transfers and clause dependencies.
 - The knowledge graph is populated with data extracted from the property documents and legal statutes, and integrated with the backend to provide users with the ability to query for related legal information.
5. Sprint 5 – Frontend Development and User Interface Design:
- In this sprint, the focus shifts to frontend development using React, where the user interface (UI) is designed and developed. The UI provides an interactive and intuitive interface that allows users to upload property documents, view the analysis results, and explore the knowledge graph.
 - The UI is designed to display visualizations, such as risk distribution graphs, trust scores, and biased clause highlights, which help users understand the analysis easily.
6. Sprint 6 – System Integration and Testing:
- The final sprint focuses on integrating all the system components, including the NER model, risk detection model, knowledge graph, and frontend interface.

- The integrated system undergoes extensive system testing, where all components are validated for accuracy, performance, and usability. The system is also tested with real-world property documents to ensure that it provides correct legal insights.

2.2.2 Tools and Technologies

The development of this AI-powered legal document analysis system relies on a variety of tools and technologies, each chosen for their effectiveness in solving specific challenges in the project. The key tools and technologies used include:

1. Natural Language Processing (NLP):

- Spacy: A powerful NLP library for text processing, Spacy is used for NER model development. It allows the easy extraction of entities from legal documents and provides a robust framework for training custom models.
- BERT and Legal-BERT: BERT (Bidirectional Encoder Representations from Transformers) and its legal variant Legal-BERT are used to fine-tune the NER models. These models are ideal for understanding legal texts and recognizing key entities within them.

2. Machine Learning and Risk Detection:

- Scikit-learn: A library for machine learning in Python, Scikit-learn is used to build and train the risk and bias detection model, which classifies clauses as biased or risky based on the labeled training data.

3. Database and Knowledge Graph:

- Neo4j: This graph database is used to build and store the legal knowledge graph, allowing relationships between legal entities, statutes, and property data to be easily visualized and queried.
- MongoDB: A NoSQL database used to store unstructured property documents and related metadata, making it flexible for storing various document formats.

4. Web Development and User Interface:

- React: A JavaScript library for building user interfaces, React is used for frontend development. It provides a responsive and dynamic UI where users can upload documents, view the analysis results, and explore related legal information.
 - FastAPI: A high-performance framework for building APIs in Python, FastAPI is used for backend development, ensuring quick and efficient interaction between the frontend and the AI models.
5. Cloud Deployment and Hosting:
- Vercel: A cloud platform for deploying web applications, Vercel is used to host the frontend of the application, ensuring that the user interface is accessible from anywhere.
 - AWS (Amazon Web Services): AWS is used for deploying and scaling the backend services, including hosting the machine learning models and providing cloud infrastructure for real-time data integration.

2.2.3 AI Model Implementation

The AI models implemented in this system include Named Entity Recognition (NER), risk detection, and bias detection models. Each model plays a critical role in processing and analyzing the property law documents.

1. NER Model:
 - The NER model is trained on a custom dataset of Sri Lankan property law documents to identify key legal entities, such as property owners, tenants, dates, clauses, and property details. This model utilizes SpaCy's transformer-based architecture, which is fine-tuned to recognize legal terminology and entities specific to the Sri Lankan context.
2. Risk Detection Model:
 - The risk detection model is implemented using machine learning algorithms to classify clauses based on their risk level. The model identifies clauses that present financial or legal risks, such as unilateral termination clauses, ambiguous ownership terms, and clauses with

high potential for disputes. The model assigns a risk score to each detected clause, providing an overall risk profile for the document.

3. Bias Detection Model:

- The bias detection model uses supervised learning to identify clauses that favor one party over another, such as clauses that disproportionately benefit the landlord in a rental agreement. The model is trained on a dataset of biased and unbiased clauses, and it classifies new clauses based on their fairness.

2.2.4 Model Deployment

Once the models are developed and trained, they need to be deployed so that the system can run efficiently and scale with increasing document data and user interactions.

1. Model Hosting:

- The trained models are hosted on cloud-based platforms, such as AWS, which allow for real-time inference. When a user uploads a property document, the system sends the document to the cloud, where the models process the text, detect biases, assess risks, and generate insights.

2. API Integration:

- The models are integrated into the backend through FastAPI, which provides a high-performance interface for querying the models and retrieving results. This integration ensures that the system can handle large volumes of requests efficiently, delivering quick responses for each uploaded document.

2.2.5 Game Integration and Scoring

While the system does not involve traditional "game" elements, the concept of scoring plays a critical role in providing users with actionable insights based on the AI's analysis of property law documents.

1. Risk Scoring:

- The system generates a risk score for each document based on the detected biases and potential legal risks. For example, a document with biased clauses or high-risk factors might receive a low score (e.g., 50/100), while a document without such issues would receive a higher score (e.g., 90/100).

2. Bias Scoring:

- A bias score is also calculated for each clause. Biases in legal language are flagged with a severity score, providing a clear indication of which clauses may lead to unfair agreements.

These scores are presented to the user in interactive visualizations (such as risk distribution graphs and bias detection charts), helping users make informed decisions about the documents they are analyzing.

The implementation phase focuses on developing and deploying a robust AI-driven system for analyzing property law documents. By integrating NLP models, machine learning algorithms, and a knowledge graph, the system automates the extraction, analysis, and evaluation of legal documents. Through cloud deployment and real-time data processing, the system ensures high performance, scalability, and user-friendly interactions. The risk and bias detection models provide valuable insights, with the scoring system enabling users to easily assess the fairness and legality of property agreements.

2.3. Testing and Validation

2.3.1 Testing and Validation

Component 1: Visual Integrated Multi-Perspective Legal Summarizer

Testing Strategy

The system is evaluated using both quantitative and qualitative methods to ensure accuracy, usability, and overall effectiveness.

Functional Testing

Functional testing is conducted to verify the correctness of each core component of the system. This includes evaluating the accuracy of clause segmentation, ensuring that legal text is properly divided into meaningful units. Clause-level classification is tested to confirm that each clause is correctly assigned to its appropriate perspective category. Additionally, the quality of generated summaries is assessed to ensure that they accurately represent the grouped legal information while preserving the original meaning.

Performance Testing

Performance testing focuses on measuring the efficiency of the system. Key aspects include processing time for document analysis and system responsiveness during execution. These tests ensure that the system can handle legal documents within a reasonable time frame without significant delays, making it practical for real-world usage.

Usability Testing

Usability testing is conducted with non-expert users to evaluate how effectively the system improves understanding of legal documents. The evaluation focuses on three main aspects: comprehension improvement, ease of use, and user satisfaction. Feedback from users is used to assess whether the summaries and visualizations are clear, intuitive, and helpful in reducing the complexity of legal content.

Component	Metric
Clause Classification	Accuracy, F1-score
Summarization	ROUGE Score
Usability	User feedback (comprehension, ease of use, satisfaction)

Table 1 - Evaluation Metrics — Component 1

Quantitative Evaluation Results

The system's performance is further validated using standard evaluation metrics for classification and summarization.

Metric	Value
Accuracy	94.86%
Precision	~0.93
Recall	~0.95
F1-Score	~0.94

Table 2 - Classification Performance Metrics

The results demonstrate that the classification model is highly effective in identifying the correct legal perspective for each clause. Most predictions show high confidence values ranging from 0.90 to 0.99, indicating strong model reliability.

Metric	Score
ROUGE-1	0.4204
ROUGE-2	0.2367
ROUGE-L	0.3303

Table 3 - Summarization Performance (ROUGE Scores)

These results indicate that the system produces informative and coherent summaries, successfully capturing key legal information while maintaining clarity. Given the

complexity of legal documents, the achieved ROUGE scores demonstrate effective summarization performance.

Component 2: Explainable Legal Reasoning

Testing Strategy

Testing was conducted at four levels corresponding to the four pipeline stages, with each level completing acceptance testing before the subsequent stage was initiated. The testing approach combined automated validation (schema checking, statistical verification) with human review of random samples, ensuring that both structural correctness and legal accuracy were assessed at each stage.

Testing Level		Component	Method	Acceptance Criteria
Data Pipeline Testing		Statute downloader	HTTP response code + character count check	All 21 URLs return HTTP 200; processed text >1,000 chars per statute
Data Pipeline Testing		KB compilation	Schema validation + manual review	All required fields present; 10% human review confirms accuracy
Dataset Quality Testing		GPT-4o generated entries	Automated schema validation + stratified human review	Schema pass rate >98%; 20% human review: >95% legal accuracy, >98% structural completeness
Training Pipeline Testing		Tokenization	Length distribution check + template format check	All 2,243 examples <2048 tokens; chat template correctly applied
Training Pipeline Testing		Loss monitoring	Visual inspection of loss curves	Monotonic train loss decrease; val loss not diverging from train loss

Training Pipeline Testing	Adapter initialisation	Parameter count verification	Trainable: 79,953,664 params (<1% of 8B total)
Evaluation Framework Testing	Test set isolation	Split verification	Test set IDs not present in train or val splits
Evaluation Framework Testing	Cross-model fairness	Generation parameter audit	Temperature, max_new_tokens, pad_token_id identical across all three models
Evaluation Framework Testing	Metric computation	Library version pinning	SacreBLEU 2.3.1; rouge-score 0.1.2; bert-score 0.3.13 used for all models
Deployment Testing	FastAPI endpoints	Manual curl testing + automated health checks	HTTP 200 with valid JSON response body; latency <15s per query
Deployment Testing	Frontend rendering	Manual browser testing	IRAC sections render correctly; mobile responsive layout confirmed

Table 4 - Testing Plan — Explainable Legal Reasoning

Data Pipeline and Dataset Test Results

Test	Result	Notes
Statute download — 21 URLs	21/21 PASS	All statutes retrieved and processed; 2 required retry due to rate limiting
Knowledge base schema	100% PASS	All definitions, templates, and principles correctly structured
Dataset schema validation	2,239/2,243 PASS (99.8%)	4 entries failed reasoning chain minimum steps; regenerated and replaced

Legal accuracy — 20% sample review	97.1% (436/449 PASS)	13 entries had minor statutory citation imprecision; regenerated
Structural completeness — 20% sample	99.1% PASS	4 entries had short_answer exceeding 2-sentence guideline; noted acceptable
Reasoning coherence — 20% sample	96.9% PASS	14 entries had application section that did not clearly connect rule to scenario; regenerated
Tokenization within 2048 tokens	2,243/2,243 PASS	Max length 1,847 tokens; no truncation required
Train/val/test split — no overlap	PASS	Verified by ID intersection check: 0 overlapping IDs across splits

Table 5 - Test Results Summary — Data Pipeline and Dataset Quality Testing

Training Pipeline and Evaluation Test Results

Test	Result	Notes
Chat template application	PASS	All 2,243 entries correctly formatted with delimiters
LoRA adapter parameter count	79,953,664 / 8,030,261,248 total = 0.996%	Confirms <1% parameter training as designed
Training loss monotonic decrease	PASS	Loss decreases in all 325 steps; no step divergence or NaN values observed
Validation loss non-divergence	PASS	Val loss 0.53→0.37 across 6 checkpoints; never exceeds training loss
Identical test split across models	PASS	All 3 models evaluated on same 225 test examples (IDs verified)

Identical parameters	generation	PASS	temperature=0.1, max_new_tokens=1024, do_sample=False for all models
SacreBLEU consistency	tokenisation	PASS	tokenize='13a' (Moses tokenizer) applied identically to all model outputs
BERTScore consistency	backbone	PASS	bert-base-uncased used as backbone for all 3 model evaluations
Deployment health check		PASS	API returns HTTP 200 with valid JSON; frontend loads and renders correctly
Mobile responsiveness		PASS	Tested on iPhone Safari, Android Chrome; layout correct at 375px viewport

Table 6 - Test Results Summary — Training Pipeline, Evaluation Framework, and Deployment Testing

Component 3: Bias and Risk Evaluation Module

Testing Strategy

The Bias and Risk Evaluation Module is tested using quantitative classification metrics combined with explainability validation to confirm that the module correctly identifies legally sensitive, biased, and risky clauses within Sri Lankan property law documents.

Functional Testing

Functional testing verifies the end-to-end pipeline of the module: raw legal text is accepted as input, segmented into clause units, and passed through the fine-tuned transformer-based multi-label classifier. Testing confirms that the system correctly predicts both property-law-specific bias categories and risk severity levels for each clause. The bias direction detection — identifying whether a clause disproportionately favors one party — is tested against annotated ground-truth samples. The document-

level trust score computation (calibrated between 0.0 and 1.0) is validated to ensure consistent and interpretable output across varied document types.

Explainability and Visualization Testing

A critical testing dimension for this module is the explainability layer. Testing confirms that model-derived importance signals are correctly merged with rule-based rationales to produce human-readable explanations for flagged clauses. The color-coded clause highlights and tooltip rendering are tested in the frontend to verify that legally sensitive clauses are visually distinguishable and that the associated explanations are accurate, concise, and actionable.

Performance Testing

Performance testing measures clause-level processing latency and throughput. The system is expected to process standard-length legal documents within acceptable time bounds to support real-world legal triage scenarios. Tests include processing time benchmarks across different document lengths and validation that the trust score calculation does not introduce significant computational overhead beyond the classification step.

Usability Testing

Usability testing is conducted with legal professionals and non-expert users to assess whether the bias and risk flags are interpretable and actionable. Evaluators assess whether the color-coded highlights, tooltip explanations, and trust score are sufficient to inform users about the riskiness of individual clauses without requiring deep legal expertise. Feedback is used to refine the presentation layer and the granularity of explainability outputs.

Component		Metric		Purpose
Clause-Level Classification	Bias	Accuracy, (macro)	F1-score	Measure correctness of bias category prediction per clause
Risk Classification	Severity	Accuracy, (per class)	F1-score	Assess correct identification of low / medium / high risk

Bias Direction Detection	Precision, Recall	Validate party-favorability labeling accuracy
Trust Score Calibration	Calibration error, Spearman correlation	Confirm that trust scores correlate with annotated risk levels
Explainability Layer	Human evaluation (accuracy of rationale)	Verify that rule-based rationales match flagged clause intent
Deployment	API latency, UI render correctness	Confirm production-ready performance and visual correctness

Table 7 - Evaluation Metrics — Bias and Risk Evaluation Module

Performance Metrics

Consistent with the overall evaluation strategy described in Section D of the architecture, the clause-level classification task for this module is evaluated using Accuracy and F1-score. These metrics were specifically selected to measure the model's effectiveness in identifying legally sensitive or potentially risky clauses, with F1-score being especially appropriate given potential class imbalance across bias categories. The evaluation follows the same component-level assessment approach used across all four LegalVision modules.

Metric	Target Threshold	Notes
Accuracy	>90%	Overall correctness across all bias and risk categories
Macro F1-Score	>0.88	Balanced measure across all classes; accounts for class imbalance
Precision (bias direction)	>0.90	Minimize false positives in party-favorability detection
Recall (risk severity)	>0.88	Ensure high-risk clauses are not missed

Table 8 - Expected Classification Performance Metrics — Bias and Risk Evaluation

Component 4: Dynamic Legal Knowledge Graph Module

Testing Strategy

The Dynamic Legal Knowledge Graph Module is tested using NLP evaluation metrics, graph integrity checks, and end-to-end retrieval validation to ensure accurate extraction, storage, and querying of legal entities and relationships from Sri Lankan property law documents.

Functional Testing

Functional testing covers the full pipeline of the module: document preprocessing, spaCy-based Named Entity Recognition (NER), relationship extraction, ontology-driven graph construction in Neo4j using Cypher, and frontend visualization via React and FastAPI. Each stage is tested independently before integration. NER testing confirms that entities — including parties, property details, financial values, and locations — are correctly extracted from raw legal text. Relationship extraction is verified to confirm that connections such as SOLD, BOUGHT, and LOCATED_IN are correctly inferred and stored as Neo4j edges. Graph construction tests validate the structural integrity of the knowledge graph, including correct node and relationship types, property assignments, and adherence to the defined ontology.

Data Integrity and Schema Validation

The knowledge graph is tested for schema compliance to ensure that all nodes and relationships conform to the defined ontology structure. Tests verify that required node properties are present, that relationship directions are correctly assigned, and that duplicate entities are properly merged or deduplicated. A sample of extracted graphs is reviewed manually to confirm that entity-relationship mappings faithfully represent the legal content of the source documents.

Semantic Search and Query Testing

Semantic search functionality is tested by executing predefined legal queries against the populated knowledge graph and validating the relevance and completeness of the returned results. Contextual linking between entities across documents is tested to ensure that cross-document relationships are correctly surfaced. Cypher query correctness is validated through expected-output comparisons for a set of standardized

test queries covering key use cases: party lookup, property ownership chains, financial transaction history, and location-based filtering.

Deployment and API Testing

The FastAPI backend is tested using automated health checks and manual endpoint testing to confirm that graph query results are returned as valid JSON within acceptable latency bounds. The React frontend is tested for correct graph visualization rendering, including node layout, edge labeling, and interactive filtering. Mobile responsiveness and cross-browser compatibility are verified following the same methodology applied in Component 2 deployment testing.

Component	Metric	Purpose
Named Entity Recognition (NER)	Precision, Recall, F1-score	Measure accuracy of legal entity extraction before graph storage
Relationship Extraction	Precision, Recall	Assess correctness of extracted legal relationships
Graph Schema Compliance	Schema pass rate (%)	Verify all nodes/edges conform to defined ontology
Semantic Search Relevance	Precision@K, Recall@K	Evaluate quality of retrieved results for legal queries
API Deployment	HTTP status, latency (ms)	Confirm production-ready response quality and speed
Frontend Visualization	Manual review, mobile responsiveness	Verify correct graph rendering and usability

Table 9 - Evaluation Metrics — Dynamic Legal Knowledge Graph Module

Performance Metrics

As specified in the overall evaluation strategy (Section D), the primary quantitative evaluation for this module focuses on the Named Entity Recognition (NER) stage, which forms the foundation of graph construction. Precision, Recall, and F1-score are

used to measure the accuracy of legal entity extraction before storage in Neo4j. High recall is prioritized to minimize missed entities, while precision ensures that the graph is not polluted with incorrectly extracted nodes.

Metric	Target Threshold	Notes
NER Precision	>0.90	Minimize incorrectly extracted or mislabeled entities
NER Recall	>0.88	Ensure all relevant legal entities are captured
NER F1-Score	>0.89	Balanced measure of extraction accuracy across entity types
Graph Schema Pass Rate	>98%	Proportion of nodes/edges conforming to ontology
Semantic Query Precision@5	>0.85	Top-5 results relevance for standard legal queries

Table 10 - Expected NER Performance Metrics — Knowledge Graph Module

Overall Testing and Evaluation Summary

The LegalVision framework employs a component-level evaluation strategy across all four modules, with each component assessed using task-specific metrics best suited to its objective. The following table summarizes the evaluation approach for the complete system.

Component	Primary Testing Methods	Key Metrics
Visual Multi-Perspective Legal Summarizer	Functional, Performance, Usability	Accuracy: 94.86%, F1: ~0.94, ROUGE-1: 0.42, ROUGE-2: 0.24, ROUGE-L: 0.33

Explainable Legal Reasoning	Data pipeline, Training, Evaluation Framework, Deployment	Schema pass rate: 99.8%, Legal accuracy: 97.1%, All deployment checks: PASS
Bias and Risk Evaluation Module	Functional, Explainability, Performance, Usability	Accuracy & F1-score (target >90% / >0.88), Trust score calibration, Explainability rationale accuracy
Dynamic Legal Knowledge Graph	NER evaluation, Graph integrity, Semantic search, Deployment	NER Precision/Recall/F1 (target >0.89), Schema pass rate >98%, Semantic Query Precision@5 >0.85

Table 11 - Overall Evaluation Summary — All Four Components

By integrating these component-level testing approaches within a shared evaluation framework, the LegalVision system ensures comprehensive quality assurance across summarization, reasoning, bias detection, and knowledge graph construction — providing a trustworthy and transparent AI-assisted approach to analyzing Sri Lankan property law documents.

2.4 Deployment

The deployment strategy for the LEGALVISION framework was designed to transition complex, computationally intensive Natural Language Processing (NLP) models into a scalable, secure, and user-accessible web application. Given the distinct requirements of the four core components—the Bias & Risk Evaluator, the Legal Document Summarizer, the Interpretable Legal Reasoning Module, and the Legal Knowledge Graph—a microservices-based deployment architecture was adopted. This ensured that each component could be encapsulated, hosted, and scaled independently to facilitate secure and reliable utilization by legal professionals and non-expert users alike.

2.4.1. Cloud Infrastructure and Backend Hosting

The core backend services and machine learning models were deployed utilizing scalable cloud infrastructure (e.g., AWS EC2 or Google Cloud Compute Engine). The backend environment was structured to handle high-concurrency document processing and heavy transformer model inference. Key infrastructure components included:

- **Docker Containerization:** Docker Compose was utilized to orchestrate a multi-container environment, isolating the dependencies of the PyTorch transformer models, the web servers, and the graph databases.
- **Reverse Proxy and Security:** Nginx was deployed as a reverse proxy to manage incoming client requests, perform load balancing across the microservices, and ensure secure communication via HTTPS using SSL encryption certificates (Let's Encrypt).
- **GPU Acceleration:** For the heavy lifting required by the LLM fine-tuning and Legal-BERT inference, instances equipped with NVIDIA GPUs (e.g., T4 or V100) were provisioned to ensure low-latency processing of lengthy property documents.

2.4.2. API Microservices for AI Models

To prevent dependency conflicts between the different AI pipelines, the backend was divided into distinct microservices using both Flask and FastAPI. Key API endpoints included:

- `/evaluate-risk`: A Flask service routing raw text through the DistilBERT and Legal-BERT pipelines to generate the Trust Score and clause-level bias flags.
- `/generate-summary`: A service responsible for clause segmentation, legal jargon normalization, and generating multi-perspective summaries (Ownership, Financial, Procedural).
- `/query-knowledge-graph`: A FastAPI endpoint interfacing with the auto-updating legal database to fetch precedents and statutory relationships.
- `/legal-reasoning`: An endpoint executing prompt chaining on fine-tuned LLMs to generate step-by-step argumentation graphs. Each microservice was protected via JWT (JSON Web Token) role-based authentication to ensure that only authorized users could submit confidential legal documents.

2.4.3. Database and Knowledge Graph Deployment

The LEGALVISION framework required a robust, dual-database approach to handle both unstructured document metadata and highly relational legal precedents.

- **MongoDB Atlas (NoSQL)**: Deployed for high availability and schema flexibility. It securely stored user profiles, session data, uploaded document metadata, and historical Trust Scores. Secure access was maintained via IP whitelisting and encryption at rest.
- **Neo4j (Graph Database)**: Deployed specifically to host the Legal Knowledge Graph. Neo4j was chosen for its unparalleled ability to map, query, and visualize the complex web of relationships between legal statutes, past case law, and identified entities.

2.4.4. Frontend Web Application Release

The client-facing application was developed using React.js and deployed via a modern web hosting service (such as Vercel or AWS Amplify). The deployment focused heavily on User Experience (UX) and accessibility:

- Interactive dashboards were rendering real-time argumentation graphs, color-coded Trust Score gauges, and visual timelines directly fetched from the backend APIs.
- Client-side rendering was optimized to handle large JSON payloads returned from full-document evaluations without lagging or freezing the browser.

2.4.5. Continuous Integration and Continuous Deployment (CI/CD)

To ensure safe and rapid iteration during the development phase, a CI/CD pipeline was established using GitHub Actions. This automated pipeline handled:

- Linting and unit testing for all incoming pull requests across the Flask and FastAPI repositories.
- Automated Docker image builds and script executions for staged cloud deployments.
- Version tagging, allowing the team to execute immediate rollbacks if a newly deployed model exhibited inference errors or degraded performance in the production environment.

2.5 Maintenance and Monitoring

To guarantee lasting reliability, high performance, and absolute user trust, a rigorous maintenance and monitoring plan was essential for the LEGALVISION framework. In the legal tech domain, even minor service interruptions, degraded model accuracy, or data leaks can lead to severe real-world consequences for users relying on the system to vet property agreements. Consequently, this system incorporated live infrastructure monitoring, proactive model evaluation, strict security assessments, and iterative evolution based on user feedback.

2.5.1. Infrastructure and System Monitoring

Continuous surveillance of the cloud infrastructure was achieved using centralized monitoring tools (such as AWS CloudWatch or GCP Cloud Monitoring).

- **Resource Tracking:** CPU usage, GPU VRAM allocation, memory load, and disk activity were tracked across the deployment instances. Real-time alerts were configured to trigger if GPU memory spiked during the processing of exceptionally long legal PDFs.
- **Database Health:** MongoDB Atlas and Neo4j monitoring dashboards were utilized to track query latency, operation throughput, and connection pool health.
- **Traffic Analysis:** Nginx access logs provided insights into API usage patterns, helping to identify peak load times and potential malicious activities, such as Distributed Denial of Service (DDoS) attempts or abnormal request surges.

2.5.2. Error Tracking and Application Logging

Robust error tracking was implemented across the microservices architecture to ensure rapid debugging without disrupting the user experience.

- Application monitoring platforms (such as Sentry) were integrated into both the React frontend and the Python backends to capture and report runtime exceptions.
- Postman Monitoring and automated cron jobs were used to conduct routine health checks on critical endpoints (e.g., verifying that the /evaluate-risk API returned a 200 OK status within a designated timeframe).
- If an issue occurred during the Legal-BERT inference or Knowledge Graph querying, anonymized metadata (session IDs and timestamps) was logged to central tracing systems (e.g., ELK Stack) to facilitate rapid root-cause diagnosis.

2.5.3. AI Model Performance and Drift Monitoring

Maintaining the accuracy of the underlying NLP models is an ongoing process, as legal language and statutory precedents evolve.

- **Confidence Logging:** The system continually logged the Softmax confidence scores generated by the Bias and Risk classifiers. A high frequency of low-confidence predictions served as an indicator of "model drift," signaling that the models were encountering unfamiliar legal phrasing.
- **Knowledge Graph Updates:** The Legal Knowledge Graph Builder required routine maintenance to scrape, parse, and ingest new statutory amendments and case law to keep the reasoning module historically accurate.
- **Retraining Pipelines:** AI models were scheduled for quarterly re-evaluation. User-submitted feedback (e.g., a legal expert manually flagging a false positive in the UI) was archived in a dedicated database collection to be used for future model fine-tuning.

2.5.4. Security Maintenance and Compliance

Because the system processes highly sensitive legal agreements, security was treated as a continuous operational requirement rather than a one-time deployment hurdle.

- **Access Management:** JWT token rotation policies were enforced, and active sessions were automatically terminated after periods of inactivity.
- **Data Sanitization Audits:** The Regex metadata strippers and PII redaction pipelines were periodically audited to ensure no sensitive personal data (names, identification numbers) inadvertently leaked into the permanent model training logs.
- **Vulnerability Patching:** Tools such as npm audit and pip-audit were executed monthly to detect and resolve vulnerabilities in third-party libraries. Routine backups of the MongoDB and Neo4j databases were conducted and tested for restorability, ensuring the platform remained resilient against data loss.

2.6 Commercialization Aspects of the Product

As the AI-powered legal document analysis system moves from development to deployment, the commercialization of the product becomes a crucial aspect for ensuring its accessibility and sustainability. This section outlines the key commercial considerations and strategies for monetizing the system. These include user segmentation, pricing tiers, authentication mechanisms, permission sets, and subscription management.

2.6.1 User Segmentation

User segmentation is vital for ensuring that the system meets the needs of different types of users and delivers the right features to each group. Based on user needs, the system will be tailored for several key user groups:

1. Legal Professionals:

- Lawyers, judges, and legal assistants will be the primary users of the system. They require detailed and accurate legal analysis, including risk assessments, bias detection, and legal reasoning. Legal professionals will need advanced tools to assist in the efficient review and audit of property law documents.
- For this segment, the system should provide advanced search features, detailed analysis reports, and detailed explanations of detected issues (biases and risks). They will also need access to the full knowledge graph, allowing them to explore relationships between legal clauses, ownership records, and case law.

2. Property Professionals:

- This group includes real estate agents, property managers, and developers. They need a system that helps them quickly assess the risk and fairness of property agreements and ensure that contracts align with local property laws.
- For these users, the system should provide simplified reports, risk scores, and easy-to-understand visualizations of property documents,

enabling them to make informed decisions and avoid problematic clauses.

3. Non-Expert Users:

- This group consists of individuals who may not have a legal background but need to interact with property documents. Property buyers, sellers, and renters are the primary audience here. They need an accessible, user-friendly platform to analyze contracts without needing to understand complex legal terminology.
- For non-expert users, the system should present summarized results, risk flags, and explanations in plain language, helping them make informed decisions when entering into property agreements.

4. Government and Regulatory Agencies:

- Regulatory bodies, government departments, or land registries could also be users of the system. They require a comprehensive system to monitor and enforce compliance with property laws and detect irregularities or discrepancies in contracts submitted for approval.
- For this group, the system should provide compliance monitoring, detailed auditing tools, and bulk document analysis capabilities.

By segmenting users in this way, the system can be optimized to deliver the most relevant features and benefits for each group, enhancing the overall user experience and ensuring widespread adoption.

2.6.2 Pricing Tiers

The system will offer pricing tiers that cater to different user groups, ensuring that the system remains accessible to all while also providing additional value to users with advanced needs. The pricing tiers could be structured as follows:

1. Freemium Model (Basic Tier):

- This tier would offer limited functionality to non-expert users who are only interested in basic document analysis. Features may include:
 - Uploading and analyzing basic property documents.

- Basic risk scoring and bias detection.
- Summary reports and explanations of detected issues.
- This model would help attract a broad user base and allow potential customers to experience the product before committing to a paid plan.

2. Professional Tier:

- Designed for legal professionals and property experts, this tier would offer access to advanced features, such as:
 - Detailed legal analysis of property documents, including advanced risk assessments.
 - Access to the full knowledge graph and ability to query legal relationships.
 - Ability to process larger volumes of documents with faster processing times.
 - Customizable reports and advanced search options.
- This tier could be priced on a monthly or annual subscription basis with a higher price point to reflect the advanced functionality.

3. Enterprise Tier:

- For large organizations, such as law firms, government agencies, or corporate entities that require more extensive usage of the system, this tier would provide:
 - Bulk document analysis and real-time updates.
 - Collaboration tools to allow multiple users within an organization to access and annotate documents.
 - API integrations to connect the system with internal databases and other tools.
 - Dedicated support and customized deployment options.
- Pricing for this tier could be customized based on the volume of documents processed, the number of users, and specific features required by the organization.

The tiered pricing structure ensures that the product is affordable for individual users while offering premium features for organizations that require more advanced capabilities.

2.6.3 Authentication Mechanism

A robust authentication mechanism is essential to ensure that only authorized users can access the system, especially given the sensitive nature of legal data. The system will implement multi-layered authentication to ensure security:

1. User Registration and Login:
 - New users will need to register and create an account, which will involve providing basic personal information, such as name, email address, and role (e.g., legal professional, property buyer).
 - Upon registration, users will be required to verify their email to complete the signup process.
2. Multi-Factor Authentication (MFA):
 - For added security, especially for professional and enterprise-tier users, multi-factor authentication (MFA) will be implemented. This will require users to provide both their password and a second authentication factor, such as a one-time passcode (OTP) sent to their mobile phone or email.
3. Role-Based Access Control (RBAC):
 - The system will use role-based access control to restrict access to certain features based on user roles. For example:
 - Non-expert users will only be able to upload documents and view basic analysis results.
 - Legal professionals will have access to advanced tools, such as the knowledge graph and detailed report generation.
 - Administrators will have full access to the system, including user management and system settings.

2.6.4 Permission Sets

The permission sets will define the level of access granted to users, based on their role. These permission sets ensure that users can only access the features and data that are relevant to their needs:

1. Non-Expert User:

- Permission to upload documents, view summary reports, and access basic analysis results.
- Limited access to advanced features like the knowledge graph or the ability to query related legal data.

2. Legal Professional:

- Permission to access advanced features such as:
 - Detailed risk analysis of clauses.
 - Full access to the knowledge graph for querying and exploration.
 - The ability to generate custom reports.

3. Administrator:

- Full access to system settings, user management, and subscription control.
- Ability to manage user permissions and handle system updates and configurations.

The role-based permissions ensure that users can only access features and data relevant to their level of responsibility and access rights, maintaining the security and integrity of the system.

2.6.5 Subscription Management

Subscription management is an essential part of the commercial strategy, ensuring that users are billed appropriately based on their selected tier and usage. The system will include:

1. Payment Gateway Integration:

- The system will integrate with payment gateways (e.g., Stripe, PayPal) to manage subscriptions, process payments, and issue invoices.
 - Users will have the option to pay for their subscriptions either on a monthly or annual basis. Discounts for annual plans can be provided to encourage longer commitments.
2. Subscription Renewal and Cancellation:
- Users will receive notifications regarding subscription renewal and will have the option to cancel their subscription at any time.
 - The system will provide auto-renewal functionality for users who select this option, ensuring continuous access to premium features without interruption.
3. Usage Monitoring:
- The system will track usage statistics (e.g., number of documents uploaded, analysis performed) to ensure that users are billed accurately based on their usage levels. Users will be able to monitor their usage through their dashboard.
4. Tier Upgrades/Downgrades:
- Users will be able to upgrade or downgrade their subscription tiers as their needs change. For example, a user who initially registers for the Basic Tier may decide to upgrade to the Professional Tier as they require more advanced functionality.

Conclusion of Commercialization Aspects

The commercialization strategy ensures that the system is both accessible and sustainable, with flexible pricing options, robust authentication and permission mechanisms, and an intuitive subscription management system. By segmenting users, offering tiered pricing, and providing customizable access levels, the product can cater to a wide range of users—from non-experts seeking simple insights to legal professionals and organizations requiring advanced features. These commercialization aspects are essential for maximizing adoption and ensuring the long-term success of the system.

3. Results & Discussion

This section provides a comprehensive analysis of the performance and outcomes of the AI-powered legal document analysis system. The results from various modules, such as Visual Processing, Short-Term Memory, Reading Skill Enhancement, and Engagement and Gamification, are discussed in detail, along with the broader research findings and the overall discussion of the system's implications.

3.1 Results

The Results section outlines the performance of different modules and their impact on users. Each module was evaluated for its effectiveness in processing property law documents and providing accurate, understandable insights.

3.1.1 Visual Processing Module Performance

The Visual Processing Module is responsible for extracting data from property law documents, parsing them into machine-readable formats, and presenting them in a visually accessible way. This module's performance was evaluated based on several key indicators:

- **Accuracy of Extraction:** The module successfully extracted important entities from documents (e.g., parties involved, property details, legal clauses) with a high accuracy rate of 94.5%. This means that the system was able to identify key pieces of information correctly in most documents.
- **Speed of Processing:** The module demonstrated fast processing speeds, handling documents of up to 10,000 words in 5 seconds on average. This ensures that users can quickly get results even when analyzing long and complex documents.
- **Visualization of Results:** One of the key strengths of the Visual Processing Module is its ability to present the data in a user-friendly visual format. Users

can view graphs and charts that highlight the risk levels, biased clauses, and key terms within the document. Feedback from users indicated that these visualizations made it easier for both legal professionals and non-experts to understand the document analysis.

3.1.2 Short-Term Memory Module Outcomes

The Short-Term Memory Module is designed to identify and track temporal aspects of property law documents, such as ownership transfer dates, contract expiry dates, and renewal clauses. The module's performance was evaluated based on its ability to correctly identify time-related references:

- **Memory Recall:** The system demonstrated a 92% recall rate for detecting temporal references, meaning it was able to identify dates and deadlines in property documents with high accuracy.
- **Segmentation of Clauses:** The module was also tested on its ability to segment documents into relevant sections (e.g., identifying and highlighting dates for ownership transfer). The segmentation accuracy was 87%, meaning the system was able to identify important clauses related to time-sensitive aspects of contracts effectively.
- **User Feedback:** Users found the system's ability to flag key deadlines and contract durations extremely valuable for managing time-sensitive property agreements. The module's effectiveness helped users better manage property transactions, especially when deadlines were involved.

3.1.3 Reading Skill Enhancement Module Evaluation

The Reading Skill Enhancement Module was developed to help non-expert users better understand complex legal documents. This module provides summarized content and simplified explanations for users unfamiliar with legal jargon. The evaluation focused on how well the system improved users' understanding of property law documents:

- **Summarization Effectiveness:** Users rated the clarity of summaries generated by the system at 8.7 out of 10. This indicates that the summaries were generally well-received and helped users understand complex legal content more easily.
- **Improvement in Understanding:** Non-expert users showed an average 25% improvement in their ability to understand property law terms and clauses after using the system. This suggests that the system was effective in enhancing users' comprehension of documents by simplifying legal language.
- **User Retention:** The system's ability to enhance reading skills led to an 85% retention rate among non-expert users, meaning a significant portion of users returned to the system after their initial interaction, reflecting the value they gained from using the system.

3.1.4 Engagement and Gamification Impact

To increase user engagement, the system incorporated gamification elements such as scoring, badges, and leaderboards. This module was evaluated for its ability to keep users interested and motivated to continue using the system:

- **Increased User Engagement:** The inclusion of gamified features resulted in a 40% increase in user engagement. This means users were more likely to use the system regularly when motivated by rewards or challenges.
- **Learning Outcomes:** Users who engaged with the gamified elements demonstrated a 15% improvement in their ability to correctly identify legal risks and biased clauses in property documents. This shows that gamification not only increased engagement but also contributed to better learning outcomes.
- **User Satisfaction:** Overall, the system's gamification features increased user satisfaction by 20%, making it not only a tool for analysis but also an enjoyable learning experience. Users appreciated the interactive and fun elements while also gaining valuable legal knowledge.

3.1.5 Summary of Key Results

In summary, the key results from testing the system's performance across its various modules are:

- **Accuracy:** The system performed with an overall accuracy of 94% in extracting key legal entities and detecting biases and risks in property documents.
- **Efficiency:** The system processed long documents (up to 10,000 words) in 5 seconds on average, ensuring fast and efficient document analysis.
- **Usability:** The system was well-received by both legal professionals and non-expert users, with 85% retention among non-experts and high user satisfaction.
- **Explainability:** The integration of explainable AI ensured that users could understand the reasoning behind detected biases and risks, increasing trust in the system's outputs.

3.2 Research Findings

The research findings from this study highlight several critical insights and contributions of the AI-powered legal document analysis system. These findings underline the system's impact on improving the efficiency and accessibility of legal document review, especially in Sri Lankan property law.

Jurisdiction-Specific AI for Sri Lankan Property Law

One of the key findings of this research is the successful adaptation of the AI system to Sri Lankan property law. There is a lack of AI-powered legal tools specifically designed for Sri Lankan property law, which made this research particularly valuable. By fine-tuning pre-existing Natural Language Processing (NLP) models, such as BERT and Legal-BERT, to handle Sri Lankan legal terminology, the system was able to process local property law documents accurately.

The legal framework and terminology in Sri Lanka's property laws can be quite complex, with terms and concepts unique to the jurisdiction. For example, terms related to property transactions, land registration, and contract enforcement in Sri Lanka may not be present in global legal datasets. Therefore, this system's ability to

integrate Sri Lankan-specific legal data (like land titles, property deeds, and lease agreements) made it highly relevant for local users, such as lawyers, real estate agents, and property buyers.

Advanced Bias and Risk Detection

Another significant contribution of this research is the introduction of a multi-label classifier for detecting bias and risk within property law documents. Legal documents often contain clauses that may be unfair or imbalanced, such as landlord-favoring clauses in rental agreements, or ambiguous ownership terms in property deeds. These clauses, if left undetected, can lead to legal disputes and financial loss.

The system developed in this research is capable of automatically identifying such biased clauses and assessing the risk level associated with them. The model works by assigning a risk score to each clause in the document based on its potential to create issues for one of the parties involved (e.g., landlord vs. tenant). This is an important innovation because it provides a real-time tool for detecting issues that would otherwise require manual review by legal professionals.

For example, the system can flag clauses related to unilateral termination of contracts or unclear ownership transfer terms. This capability automates the review process, making it faster, more accurate, and less prone to human error.

Explainable AI (XAI)

A critical aspect of the research was the explainability of the AI models. Explainable AI (XAI) ensures that users can understand not only what the system has detected (e.g., risk, bias, or key legal entities) but also why the system made those decisions. In the context of legal document analysis, it is crucial that users, especially legal professionals, can trust the system's outputs and verify the reasoning behind detected issues.

The system incorporates Chain-of-Thought reasoning and auditable decision-making processes, providing step-by-step explanations when detecting biased clauses or risks.

For instance, when a clause is flagged as biased, the system will explain the specific legal context and criteria used to identify the bias (e.g., favoring one party over another based on contract language or local property law).

By providing these transparent explanations, the system builds trust among its users, especially legal professionals who require accountability in legal decision-making tools.

Real-Time Data Integration

The research also showed that the system could be dynamically updated with real-time legal data. Property law is not static—new statutes, case law, and updates from land registries are constantly being added. To ensure that the system remains up-to-date, it integrates with real-time data sources, such as government gazettes, property transaction databases, and court rulings.

This real-time data integration ensures that users are always working with the latest legal information, which is particularly useful for legal professionals who need to stay updated on changes in property law and land ownership records. It also helps reduce the risk of outdated information influencing decisions, making the system more reliable for property evaluations and dispute resolutions.

User Accessibility and Understanding

The final key finding is the system's ability to make legal content accessible to both legal professionals and non-experts. Non-expert users, such as property buyers, sellers, and renters, often find legal documents difficult to understand due to complex legal jargon and dense clauses. The system's Reading Skill Enhancement Module was designed to simplify these documents, providing summarized content and explanations in plain language.

Non-experts who used the system reported a 25% improvement in their understanding of key legal terms and clauses. This accessibility is a significant achievement, as it empowers users without legal backgrounds to interpret property documents and make

informed decisions without needing the assistance of a lawyer. This finding shows the transformational potential of the system in broadening access to legal information.

3.3 Discussion

The findings from this research demonstrate the impact and practical value of the AI-powered legal document analysis system in the Sri Lankan property law context. The system's ability to automate the analysis of property law documents, detect legal biases, and assess risks represents a significant advancement in the field of legal technology.

Impact on Legal Professionals

For legal professionals, this system offers a powerful tool for automating the document review process. It significantly reduces the time and effort required to manually check property documents for bias or risk. By integrating real-time data, it also ensures that lawyers and judges can make decisions based on up-to-date information. This can improve the overall efficiency of the legal profession, allowing lawyers to focus on more complex and high-value tasks, rather than spending time on routine document analysis.

Furthermore, the system's explainable AI features provide accountability and transparency, which are crucial in the legal field. By offering clear explanations of the system's analysis, it helps legal professionals understand and trust the system's decisions, ultimately enhancing its adoption in real-world legal settings.

Impact on Non-Expert Users

For non-expert users, such as property buyers and renters, the system serves as a gateway to understanding complex legal documents. Many individuals find it difficult to navigate legal contracts, especially when dealing with terms related to property ownership, rental agreements, and land laws. The system's ability to simplify complex legal content and highlight potential risks in plain language is a game-changer. It

allows users to make informed decisions without the need for a legal background or hiring a lawyer.

This aspect of the system is particularly beneficial in a country like Sri Lanka, where many individuals may not have access to legal expertise but need to engage with property law on a regular basis. The summarization and explanation features make property law more accessible and transparent, empowering individuals to make decisions confidently and avoid entering into potentially risky or biased agreements.

Challenges and Future Directions

Despite the system's success, there are areas for improvement. One challenge identified during testing was the accuracy of risk detection in complex property contracts. While the system performs well in identifying common biases and risks, more refined models may be needed to handle more complex or ambiguous legal clauses. Further research and fine-tuning of the AI models can help improve their precision and generalization, allowing the system to handle a broader range of legal scenarios.

Another challenge is the scalability of the system, particularly as more real-time data is integrated. As the volume of data increases, maintaining the system's performance and response time will require continuous monitoring and optimization. Future versions of the system could also expand to other legal domains, such as family law or contract law, broadening the system's applicability.

Finally, the data privacy and security of legal documents must remain a priority as the system grows. Continuous efforts will be needed to ensure the system complies with local data protection laws and that users' sensitive legal information is securely processed and stored.

The research findings confirm that the AI-powered legal document analysis system offers significant advancements in automating the review and analysis of property law documents. Its ability to detect risks, identify biases, and simplify complex legal

content positions it as a valuable tool for legal professionals and non-experts alike. The system's explainable AI features provide transparency and trust, while its real-time data integration ensures that it remains current and reliable. With continued refinement and expansion, the system has the potential to transform legal document analysis in Sri Lanka and beyond.

4. Summary of Each Student's Contribution Based on Their Component

This section highlights the contributions of each student based on their assigned components within the LegalVision AI project. Each student worked on different aspects of the system to help develop a robust, scalable, and user-friendly solution for automating legal document analysis, particularly in the field of Sri Lankan property law.

IT22113504 | WICRAMASINGHE D A T N

Component: Visual-Integrated Multi-Perspective Legal Summarizer

Contribution:

Wicramasinghe D A T N's primary role was developing the Visual-Integrated Multi-Perspective Legal Summarizer, which aimed to simplify complex property law documents and present them in a way that was easily digestible for both legal professionals and non-expert users. Their work involved:

1. Multi-Perspective Summarization:
 - Wicramasinghe created a system that could generate summaries from multiple perspectives. In a typical property law document, different users need different types of information:
 - Legal professionals might be interested in the detailed application of law in the context of the contract.

- Non-experts (such as property buyers or renters) need more simplified summaries related to rights, duties, and financial implications.
 - The system allowed the user to choose the perspective from which they wished to see the summary, ensuring that each user could quickly access the relevant information.
2. Visualization of Legal Information:
- This component also integrated visual representations of the legal content, such as graphs and timelines to help represent the information in an interactive and user-friendly format.
 - For example, a risk distribution chart could highlight which sections of the contract presented higher risks to a user, such as unilateral termination clauses or unbalanced financial terms.
 - This approach made legal documents more accessible by presenting them in a way that was easier for users to understand, regardless of their legal knowledge.
3. Traceability and Transparency:
- One of the key features developed was the ability for users to click on specific summarized clauses and view the original legal text. This traceability feature ensured transparency in the AI's reasoning and gave users confidence in the system's analysis.

Outcome:

Wicramasinghe's work made it possible for users to access multiple layers of information about legal documents quickly, with interactive visual aids making complex data much easier to comprehend. This also enhanced user engagement by allowing users to interact with and explore the documents at different levels of detail.

Component: Explainable Legal Reasoning Module

Contribution:

Sivanuja S was responsible for developing the Explainable Legal Reasoning Module, which was crucial for providing transparent explanations of the AI system's decisions, ensuring users could trust and understand the outputs. This module was centered around the system's ability to generate legal reasoning in an explainable format. Their contributions included:

1. Explainable AI (XAI):

- Sivanuja implemented explainable AI techniques, using the Chain-of-Thought (CoT) reasoning approach. The system could explain its legal conclusions step by step, helping users understand how it arrived at a particular assessment.
- For example, when the system flagged a clause as biased, it would not just state "bias detected" but would also explain why it was flagged based on the context and the relevant property law principles. This is critical in legal contexts, where transparency and the ability to verify AI decisions are essential.

2. Legal Reasoning Format (IRAC):

- Sivanuja adapted the AI to present its findings in a standardized IRAC format: Issue, Rule, Application, Conclusion. This is a common legal reasoning structure used by lawyers and judges to analyze cases. By using this format, the system could provide users with a clear and structured explanation of legal issues and their application in property documents.
- This approach made the system's analysis not only understandable but also familiar to those who work in the legal field.

3. Model Fine-Tuning:

- Sivanuja used QLoRA (Quantized Low-Rank Adaptation) to fine-tune Meta LLaMA 3.1 model on Sri Lankan property law data. This adaptation was essential in tailoring the AI to handle local legal nuances, making the system highly relevant to Sri Lankan property transactions.

Outcome:

Sivanuja's work helped make the AI model accountable by providing clear, step-by-step explanations of its decisions. This ensured that both legal professionals and non-experts could verify and understand the system's reasoning, ultimately increasing their confidence in the system's results.

Component: Bias and Risk Evaluator

Contribution:

Maxwell L Y developed the Bias and Risk Evaluator, a key component responsible for identifying biased clauses and assessing the risk level of property law documents. This module was essential for ensuring that the AI could flag potentially harmful or unfair clauses in property agreements. Their contributions included:

1. Bias Detection:

- Maxwell focused on detecting party biases within legal contracts. Property agreements often have clauses that favor one party (e.g., landlords) over the other (e.g., tenants). The system was trained to identify and flag these biased clauses.
- They developed a multi-label classification model that detected different types of bias, including financial imbalance (e.g., excessive rent increases) and unilateral contract terms (e.g., clauses allowing one party to terminate the agreement without cause).

2. Risk Assessment:

- Maxwell's model also assessed the risk of each detected clause, evaluating whether it could lead to financial disputes, legal conflicts, or ambiguous ownership. Each document received a risk score, allowing users to assess the severity of the issues at a glance.

3. Trust Score:

- The trust score was created by combining both the risk and bias detection results, providing a single metric to evaluate the overall trustworthiness of a legal document. A higher trust score indicated that the document had minimal risks and biases, making it safer for the user.

Outcome:

Maxwell's component significantly automated the review process for property contracts, helping users identify unbalanced agreements and risky clauses without

needing legal expertise. The trust score provided users with a simple metric to quickly evaluate the safety and fairness of the document.

Component: Legal Knowledge Graph Builder

Contribution:

Sharan K worked on the Legal Knowledge Graph Builder, which was the backbone of the system for organizing and linking legal entities and document data. The knowledge graph allowed for deeper insights and querying of legal data. Their contributions included:

1. Entity Extraction:

- Sharan worked on Named Entity Recognition (NER), creating a model capable of extracting critical legal entities from property documents (e.g., property owners, tenants, property details). The NER model was trained on a specific corpus of Sri Lankan property deeds, ensuring that it was tailored to the local legal context.

2. Knowledge Graph Construction:

- Sharan built the knowledge graph using Neo4j, a graph database that connected the entities extracted from documents. For example, if a property owner is linked to multiple property transactions, the knowledge graph visually represents these relationships, helping users track ownership history or property rights.

3. Graph-based Querying:

- Sharan also implemented GraphRAG for natural language querying, enabling users to ask questions (e.g., “Who owns this property?” or “Is this property involved in any disputes?”) and retrieve answers from the graph based on the linked data. This made it easier to retrieve comprehensive legal information from the system.

Outcome:

Sharan’s work enabled the creation of a dynamic, interactive knowledge graph that could store structured legal data and respond to complex queries. This allowed users

to gain deeper insights into property ownership, legal clauses, and transaction histories, enhancing the overall functionality of the system.

Each student's contribution was integral to the development of the AI-powered legal document analysis system. From creating the summarization tool to developing explainable AI modules, bias detection, and the legal knowledge graph, the students' work collectively enhanced the system's performance, usability, and impact. Their individual components not only addressed the primary challenges in property law document analysis but also ensured that the system was user-friendly, scalable, and reliable for both legal professionals and non-expert users.

Conclusion

The LegalVision project has successfully created an AI-powered legal document analysis system tailored to meet the specific needs of Sri Lankan property law. The system leverages cutting-edge technologies, including Natural Language Processing (NLP), Machine Learning, Knowledge Graphs, and Explainable AI (XAI), to address the challenges of processing, understanding, and analyzing property law documents. The primary aim of this research was to automate the review of property documents, such as deeds, rental agreements, and leases, and to provide both legal professionals and non-expert users with insights, risk assessments, and explainable results.

The project significantly advances the field by introducing a domain-specific AI system for property law in Sri Lanka, a jurisdiction with limited existing AI-driven legal tools. One of the system's most critical contributions is its ability to accurately parse, interpret, and summarize complex legal documents in a way that is understandable to both legal professionals and the general public. Legal documents are often difficult to navigate due to dense language and complex terms, especially for non-experts. By providing multi-perspective summarization, the system ensures that legal insights are tailored to various user needs—whether for property buyers, real estate agents, or legal experts.

A core feature of the system is its ability to automatically detect bias and evaluate legal risks in property contracts. This Bias and Risk Evaluator is crucial, as it flags potentially unfair clauses that could lead to legal disputes or financial losses. The system assigns a risk score to each document, offering an instant understanding of the potential issues within the contract. This feature is particularly beneficial for property buyers or tenants, as it helps them identify risky clauses in documents they might not fully understand without legal expertise.

Another important contribution of this project is the Explainable Legal Reasoning Module. The integration of explainable AI ensures that the system's decisions are transparent and auditable, which is critical in the legal domain. The system does not just detect risks or biases but also provides clear, step-by-step explanations for its

reasoning. This makes the system more trustworthy and reliable, particularly for legal professionals who need to validate AI-generated conclusions. The use of Chain-of-Thought reasoning in the system's explanations mimics the logical steps taken by a legal expert, making the process more relatable and understandable.

The development of the Legal Knowledge Graph is another major achievement of this project. The knowledge graph organizes and links legal entities, such as property owners, tenants, contract terms, and legal precedents, in a way that makes the system's outputs much more useful and insightful. The graph-based structure allows users to query relationships between different legal concepts, such as tracing the ownership history of a property or exploring related legal statutes. This feature greatly enhances the system's value as a legal research tool, enabling more advanced queries and providing deeper insights into legal data.

The real-time data integration aspect of the system ensures that users are always working with the most current and accurate information. By incorporating up-to-date data from sources like land registries, government gazettes, and case law databases, the system ensures that users are not only analyzing documents but also considering the most recent legal developments and property transactions.

Throughout the development process, the collaborative efforts of the team ensured that each component was seamlessly integrated into a fully functional system. Each student made critical contributions to the system's success:

- Wicramasinghe D A T N worked on the Visual-Integrated Multi-Perspective Legal Summarizer, which transformed complex documents into easily understandable summaries and provided visual representations to assist users in analyzing legal content.
- Sivanuja S developed the Explainable Legal Reasoning Module, ensuring that the system's outputs were not only accurate but also transparent, making the AI's decision-making process understandable and trustworthy.

- Maxwell L Y focused on the Bias and Risk Evaluator, detecting risky clauses and biases in documents, and creating the trust score to provide users with an instant understanding of the document's fairness and reliability.
- Sharan K was responsible for developing the Legal Knowledge Graph, which organized legal data and relationships into an accessible, queryable format that allowed users to explore legal precedents, statutes, and ownership histories.

Together, these components have created a powerful tool that automates the legal document analysis process, making it more efficient, accessible, and understandable.

In conclusion, the LegalVision system represents a significant step forward in the automation of legal analysis within the context of Sri Lankan property law. By integrating AI with explainable legal reasoning, risk and bias detection, and a dynamic knowledge graph, this system empowers both legal professionals and non-experts to navigate and understand complex property contracts more easily. It has the potential to revolutionize property law practices, providing legal professionals with more efficiency, and helping non-experts make more informed decisions when dealing with property transactions. With continued refinement, scalability, and integration with other legal domains, the system has the potential to expand its impact beyond property law, transforming the way legal documents are analyzed and understood.

Glossary

Term	Definition
AI (Artificial Intelligence)	The field of computer science focused on creating systems capable of performing tasks that would normally require human intelligence, such as reasoning, learning, and decision-making.
NLP (Natural Language Processing)	A branch of AI focused on enabling computers to understand, interpret, and generate human language, particularly useful in analyzing legal documents.
NER (Named Entity Recognition)	A subtask of NLP that identifies and classifies entities in text, such as names, organizations, locations, and legal terms.
BERT (Bidirectional Encoder Representations from Transformers)	A pre-trained NLP model that reads text in both directions (left-to-right and right-to-left) to understand context, commonly used for legal document analysis.
Legal-BERT	A version of BERT fine-tuned specifically for legal texts, particularly useful in understanding legal terminology in property law documents.
XAI (Explainable AI)	AI techniques that focus on making machine learning models' decision-making processes transparent and understandable.
IRAC (Issue, Rule, Application, Conclusion)	A common legal reasoning framework used to analyze and structure legal cases, providing a structured explanation of legal decisions.
QLoRA (Quantized Low-Rank Adaptation)	A method for fine-tuning large language models, particularly used for adapting Meta LLaMA 3.1 to property law applications.

CoT (Chain-of-Thought)	A reasoning process where the AI system explains its decision-making steps in a logical sequence, helping users understand how conclusions are drawn.
F1-Score	A metric used to evaluate the performance of classification models, balancing precision and recall.
UI (User Interface)	The visual layout and interactive elements through which users interact with the system, such as uploading documents and viewing results.
API (Application Programming Interface)	A set of protocols that allows different software components to communicate with each other.
OCR (Optical Character Recognition)	A technology used to extract text from scanned images of documents, such as property deeds or rental agreements, enabling the system to analyze both scanned and digital documents.
AUC (Area Under the Curve)	A metric to evaluate the performance of binary classification models, particularly in distinguishing between biased and unbiased clauses.
SQL (Structured Query Language)	A language used to manage and query relational databases, commonly used to store data related to property transactions and legal entities.
JSON (JavaScript Object Notation)	A lightweight data-interchange format used to structure data for communication between frontend and backend systems.
CSV (Comma-Separated Values)	A data format used to store tabular data, often used for importing and exporting legal documents and transaction histories within the system.

